

# CloudRaid: Hunting Concurrency Bugs in the Cloud via Log-Mining

Jie Lu, Feng li, Lian li, XiaoBing Feng

State Key Laboratory of Computer Architecture, Institute of Computing Technology, Chinese Academy of Sciences

University of Chinese Academy of Sciences

August 14, 2019



# Distributed Systems

- Open Source

# Distributed Systems

- Open Source



YARN



HDFS



HBase



Cassandra

# Distributed Systems

- Open Source



YARN



HDFS



HBase



Cassandra

- Industry



Aliyun



AWS



Azure



Google cloud

# Outages

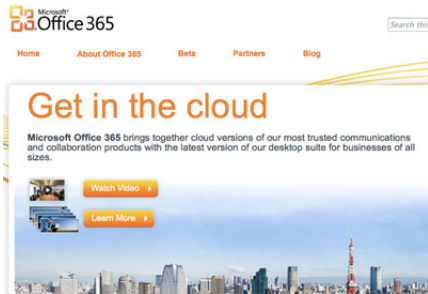
AWS, February 28, 2017



- Storage Service Disruption
- Shook the industry
- About four hours
- Wrong command

# Outages

## Microsoft Office 365, March 21, 2017



- Service become inaccessible
  - OneDrive
  - Skype
  - Outlook
- An hour
- Software bugs

# Threats of Reliability



Power Outages

---

<sup>1</sup>gunawi2016does.

# Threats of Reliability



Power Outages



Security Attacks

---

<sup>1</sup>gunawi2016does.



# Threats of Reliability



Power Outages



Security Attacks



Human Errors

---

<sup>1</sup>gunawi2016does.

# Threats of Reliability

-  Power Outages

-  Security Attacks

-  Human Errors

•  Software bugs<sup>1</sup>

---

<sup>1</sup>gunawi2016does.

# Concurrency bugs

- Local

# Concurrency bugs

- Local
  - Schedule of Threads

# Concurrency bugs

- Local
  - Schedule of Threads
  - Well studied and detected 😊

# Concurrency bugs

- Local
  - Schedule of Threads
  - Well studied and detected 😊
- Distributed

---

<sup>2</sup>leesatapornwongsa2016taxdc.

# Concurrency bugs

- Local
  - Schedule of Threads
  - Well studied and detected 😊
- Distributed
  - wrong order of messages<sup>2</sup>

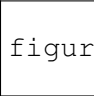
---

<sup>2</sup>leesatapornwongsa2016taxdc.

# Concurrency bugs

- Local

- Schedule of Threads
- Well studied and detected 😊



figures/right.png

- Distributed

- wrong order of messages<sup>2</sup>

---

<sup>2</sup>leesatapornwongsa2016taxdc.



# Concurrency bugs

- Local

- Schedule of Threads
- Well studied and detected 😊

figures/error.png

- Distributed

- wrong order of messages<sup>2</sup>

---

<sup>2</sup>leesatapornwongsa2016taxdc.

# Concurrency bugs

- Local

- Schedule of Threads
- Well studied and detected 😊

figures/error.png

- Distributed

- wrong order of messages<sup>2</sup>

---

<sup>2</sup>leesatapornwongsa2016taxdc.

# Concurrency bugs

- Local

- Schedule of Threads
- Well studied and detected 😊

figures/error.png

- Distributed

- wrong order of messages<sup>2</sup>
- Hard to detect 😞

---

<sup>2</sup>leesatapornwongsa2016taxdc.

# Model Checking

- Enumerating all messages order

# Model Checking

- Enumerating all messages order
  - a b c

# Model Checking

- Enumerating all messages order
  - a b c
  - (a,b,c) (a,c,b) **(b, a, c)** (b,c,a) (c, a, b) (c, b, a)

# Model Checking

- Enumerating all messages order
  - a b c
  - (a,b,c) (a,c,b) **(b, a, c)** (b,c,a) (c, a, b) (c, b, a)
- Simple input of YARN: **5,495** messages

# Model Checking

- Enumerating all messages order
  - a b c
  - (a,b,c) (a,c,b) **(b, a, c)** (b,c,a) (c, a, b) (c, b, a)
- Simple input of YARN: **5,495** messages
  - State-space explosion problem



# Model Checking

- Enumerating all messages order
  - a b c
  - (a,b,c) (a,c,b) **(b, a, c)** (b,c,a) (c, a, b) (c, b, a)
- Simple input of YARN: **5,495** messages
  - State-space explosion problem

• Are all enumerations needed?

# Model Checking

- Enumerating all messages order
  - a b c
  - (a,b,c) (a,c,b) **(b, a, c)** (b,c,a) (c, a, b) (c, b, a)
- Simple input of YARN: **5,495** messages
  - State-space explosion problem

• Are all enumerations needed?

**NO.**

# Observations from Real Bugs

- Observation 1 :

*message order  $\implies$  handler order  $\implies$  shared object access order*

---

<sup>3</sup>leesatapornwongsa2016taxdc.

# Observations from Real Bugs

- Observation 1 :

*message order  $\implies$  handler order  $\implies$  shared object access order*

- Observation 2 : Many message orders are already tested in live systems

---

<sup>3</sup>leesatapornwongsa2016taxdc.

# Observations from Real Bugs

- Observation 1 :

*message order  $\implies$  handler order  $\implies$  shared object access order*

- Observation 2 : Many message orders are already tested in live systems
- Observation 3<sup>3</sup>: Most distributed concurrency bugs can be triggered by reordering only **two** messages

---

<sup>3</sup>leesatapornwongsa2016taxdc.

# Suspicious message order

- Observation 3 :  $\langle S, P \rangle$

# Suspicious message order

- Observation 3 :  $\langle S, P \rangle$
- *Observation1*  $\Rightarrow$  *Rule1*: S and P access the same object

# Suspicious message order

- Observation 3 :  $\langle S, P \rangle$
- *Observation1*  $\Rightarrow$  *Rule1*: S and P access the same object
- *Observation2*  $\Rightarrow$  *Rule2*: The order should not be tested in live system



# Information from live system

- Non - interference in live system!

# Information from live system

- Non - interference in live system!
- Distributed systems produce massive logs for post-mortem debug!

# Information from live system

- Non - interference in live system!
- Distributed systems produce massive logs for post-mortem debug!
- Information

# Information from live system

- Non - interference in live system!
- Distributed systems produce massive logs for post-mortem debug!
- Information
  - Constant text



August 14, 2019 Start request for container\_0001

# Information from live system

- Non - interference in live system!
- Distributed systems produce massive logs for post-mortem debug!
- Information

- Constant text

August 14, 2019 Start request for container\_0001

- Time

# Information from live system

- Non - interference in live system!
- Distributed systems produce massive logs for post-mortem debug!
- Information

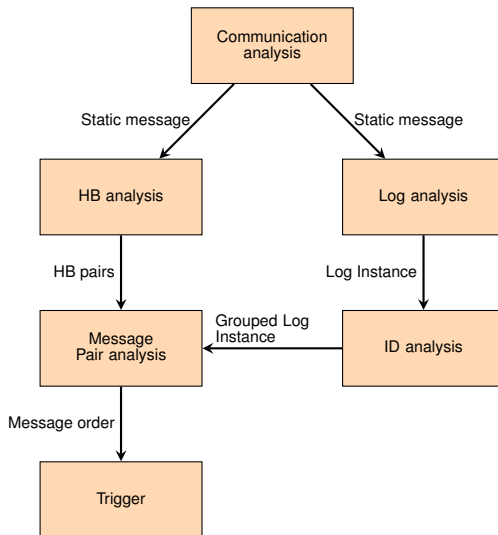
- Constant text

August 14, 2019 Start request for container\_0001

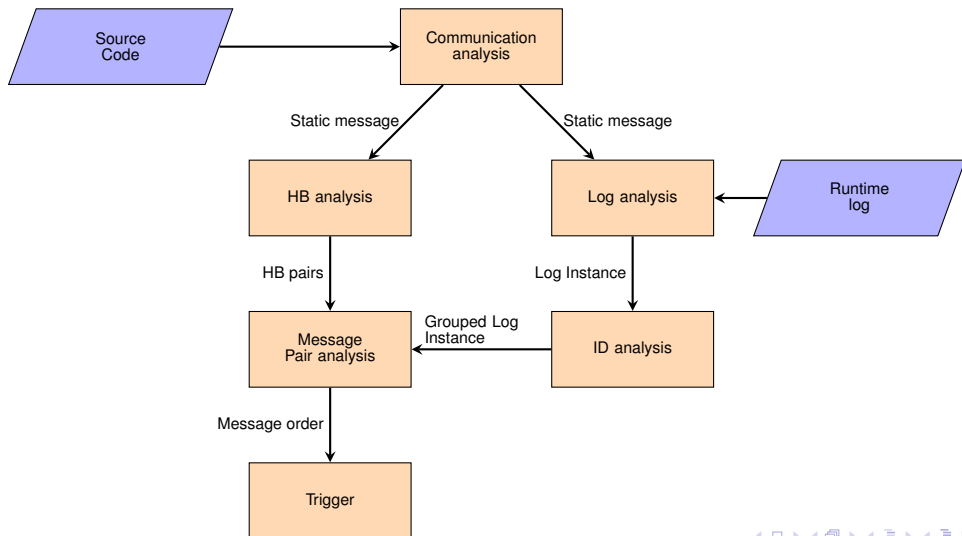
- Time

- Variable

# Overveiw of CloudRaid

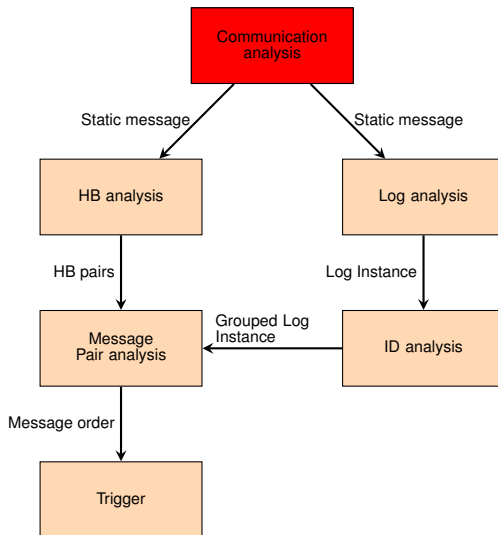


# Overveiw of CloudRaid





# Overveiw of CloudRaid



# Communication analysis

- Target messages?

# Communication analysis

- Target messages?
  - RPC, Event, Thread

# Communication analysis

- Target messages?
  - RPC, Event, Thread
  - *Client*  $\Rightarrow$  *Server*

```
//Client
public void EventProcessor.run() {
    .....
    response = proxy.startContainer(new StartConReq());
    .....
}

//Server
public StartConRes startContainer(StartConReq req) {
    ID containerID = req.getConLauContext().getConId();
    LOG.info("Start request for " + containerID);
    .....
}
```

# Communication analysis

- Target messages?
  - RPC, Event, Thread
  - *Client*  $\Rightarrow$  *Server*
  - Static analysis

```
//Client
public void EventProcessor.run() {
    .....
    response = proxy.startContainer(new StartConReq());
    .....
}

//Server
public StartConRes startContainer(StartConReq req) {
    ID containerID = req.getConLauContext().getConId();
    LOG.info("Start request for " + containerID);
    .....
}
```

# Communication analysis

- Target messages?

- RPC, Event, Thread
- *Client*  $\Rightarrow$  *Server*
- Static analysis

- Static message:  $\langle C, F, L \rangle$

```
//Client
public void EventProcessor.run() {
    .....
    response = proxy.startContainer(new StartConReq());
    .....
}

//Server
public StartConRes startContainer(StartConReq req) {
    ID containerID = req.getConLauContext().getConId();
    LOG.info("Start request for " + containerID);
    .....
}
```

# Communication analysis

- Target messages?
  - RPC, Event, Thread
  - *Client*  $\Rightarrow$  *Server*
  - Static analysis
- Static message:  $\langle C, F, L \rangle$ 
  - C: Client

```
//Client
public void EventProcessor.run() {
    .....
    response = proxy.startContainer(new StartConReq());
    .....
}

//Server
public StartConRes startContainer(StartConReq req) {
    ID containerID = req.getConLauContext().getConId();

    LOG.info("Start request for " + containerID);
    .....
}
```

# Communication analysis

- Target messages?

- RPC, Event, Thread
- *Client*  $\Rightarrow$  *Server*
- Static analysis

- Static message:  $\langle C, F, L \rangle$

- C: Client
- F: Server

```
//Client
public void EventProcessor.run() {
    .....
    response = proxy.startContainer(new StartConReq());
    .....
}

//Server
public StartConRes startContainer(StartConReq req) {
    ID containerID = req.getConLauContext().getConId();
    LOG.info("Start request for " + containerID);
    .....
}
```



# Communication analysis

- Target messages?

- RPC, Event, Thread
- *Client*  $\Rightarrow$  *Server*
- Static analysis

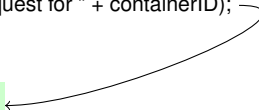
- Static message:  $\langle C, F, L \rangle$

- C: Client
- F: Server
- L: Log pattern

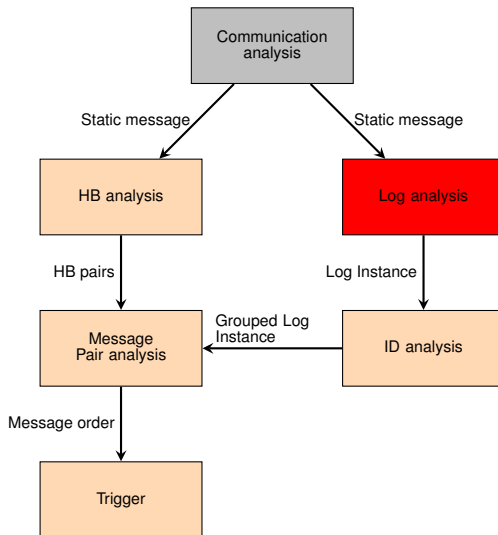
```
//Client
public void EventProcessor.run() {
    .....
    response = proxy.startContainer(new StartConReq());
    .....
}

//Server
public StartConRes startContainer(StartConReq req) {
    ID containerID = req.getConLauContext().getConId();
    LOG.info("Start request for " + containerID);
    .....
}
```

(.\*) Start request for (.\*)



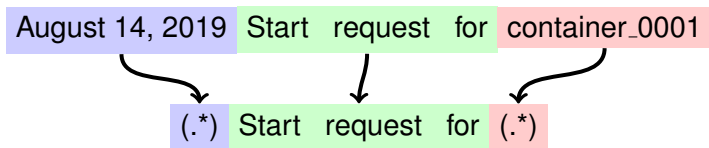
# Log analysis



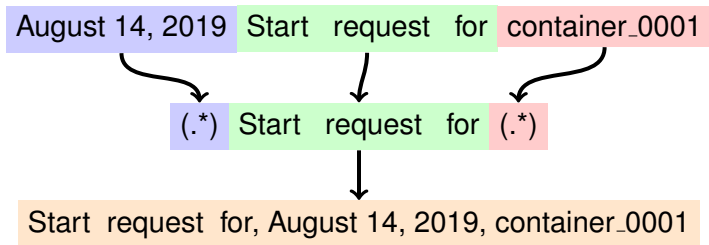
# Log analysis

August 14, 2019 Start request for container\_0001

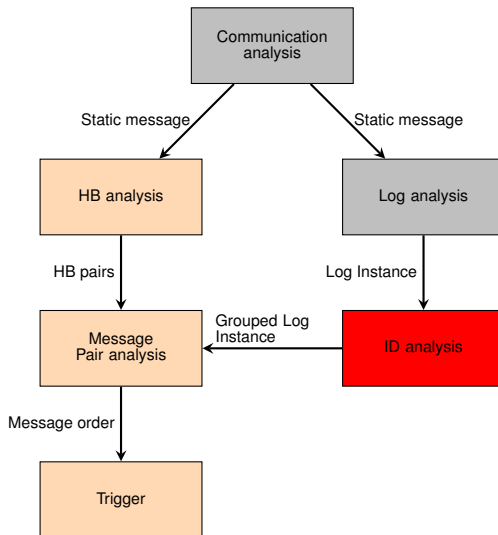
# Log analysis



# Log analysis



# ID analysis



# ID analysis

- ID is a variable

# ID analysis

- ID is a variable
- The value of one ID can index a task



# ID analysis

- ID is a variable
- The value of one ID can index a task
- What variable is ID?

```
//Server
public StartConRes startContainer(StartConReq req) {

    ID containerID = req .getConLauContext(). getConId ();

    LOG.info("Start request for " + containerID );

    .....
}
```

# ID analysis

- ID is a variable
- The value of one ID can index a task
- What variable is ID?
  - propagated from the arguments of message handler

```
//Server
public StartConRes startContainer(StartConReq req) {

    ID containerID = req .getConLauContext(). getConId ();

    LOG.info("Start request for " + containerID );

    .....
}
```

# ID analysis

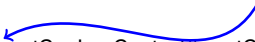
- ID is a variable
- The value of one ID can index a task
- What variable is ID?
  - propagated from the arguments of message handler

```
//Server
public StartConRes startContainer(StartConReq req) {

    ID containerID = req .getConLauContext(). getConId ();

    LOG.info("Start request for " + containerID );

    .....
}
```



# ID analysis

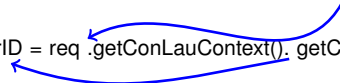
- ID is a variable
- The value of one ID can index a task
- What variable is ID?
  - propagated from the arguments of message handler

```
//Server
public StartConRes startContainer(StartConReq req) {

    ID containerID = req .getConLauContext(). getConId ();

    LOG.info("Start request for " + containerID );

    .....
}
```

A blue curved arrow originates from the 'req' parameter in the function signature and points to the 'req' object in the method call 'req.getConLauContext().getConId()'. Another blue curved arrow originates from the 'getConId()' method call and points to the 'containerID' variable in the assignment statement 'ID containerID = ...'.

# ID analysis

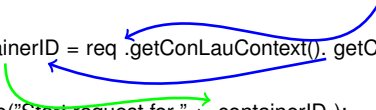
- ID is a variable
- The value of one ID can index a task
- What variable is ID?
  - propagated from the arguments of message handler
  - Printed in log

```
//Server
public StartConRes startContainer(StartConReq req) {

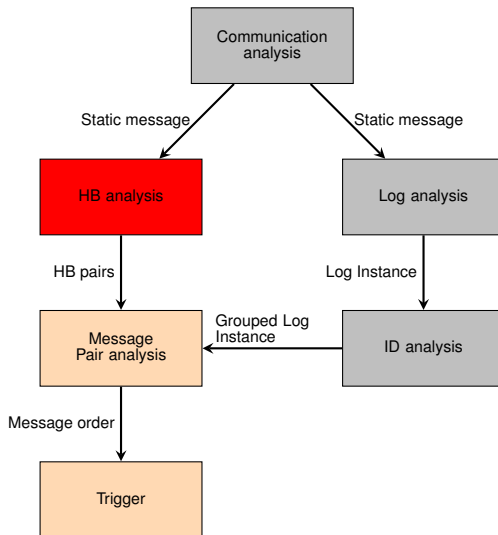
    ID containerID = req.getConLauContext().getConId();

    LOG.info("Start request for " + containerID);

    .....
}
```



# Happend Before analysis



# Happend Before analysis

- Static analysis

# Happend Before analysis

- Static analysis

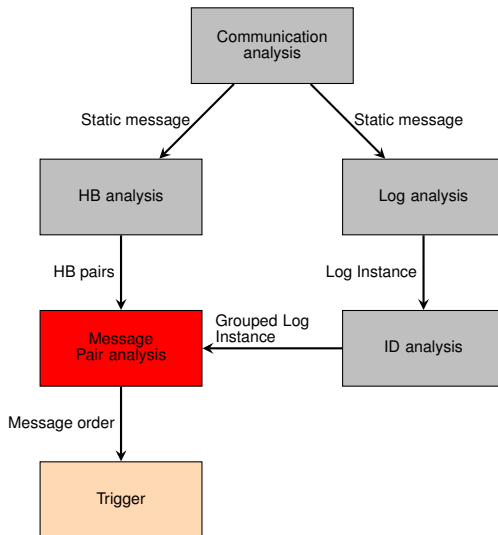
$S : \langle C_S, F_S, L_S \rangle, P : \langle C_P, F_P, L_P \rangle$

*If  $C_S$  is in  $F_P$ , then  $P$  HB  $S$*

- Fix-point



# Message Pair analysis



# Message Pair analysis

- Dynamic messages

# Message Pair analysis

- Dynamic messages
  - Static Message, RunTime logs

$S_1$   
container\_0001

$P_1$   
container\_0001

$S_2$   
container\_0002

$P_2$   
container\_0002

# Message Pair analysis

- Dynamic messages
  - Static Message, RunTime logs
  - ID value to group messages that belong to same task

$S_1$   
container\_0001


$P_1$   
container\_0001

$S_2$   
container\_0002

$P_2$   
container\_0002


# Message Pair analysis

- Dynamic messages
  - Static Message, RunTime logs
  - ID value to group messages that belong to same task
  - Time





# Message Pair analysis

- Dynamic messages
  - Static Message, RunTime logs
  - ID value to group messages that belong to same task
  - Time
- Same Shared Object(rule1) : ID value→same task




# Message Pair analysis

- Dynamic messages
  - Static Message, RunTime logs
  - ID value to group messages that belong to same task
  - Time
- Same Shared Object(rule1) : ID value→same task
- Filter executed order pairs(rule2)



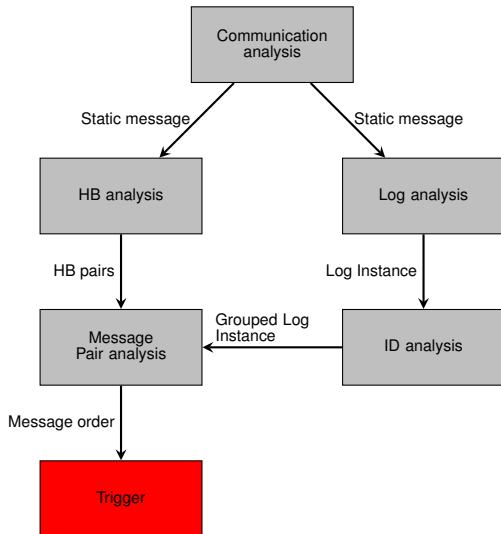
# Message Pair analysis

- Dynamic messages
  - Static Message, RunTime logs
  - ID value to group messages that belong to same task
  - Time
- Same Shared Object(rule1) : ID value  $\rightarrow$  same task
- Filter executed order pairs(rule2)

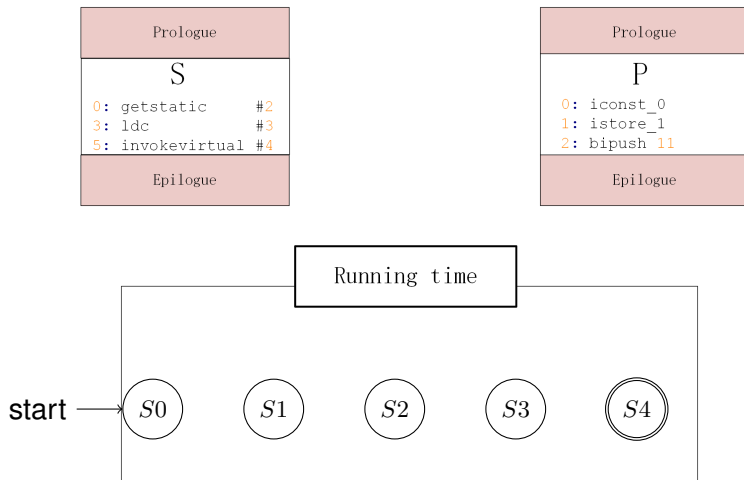




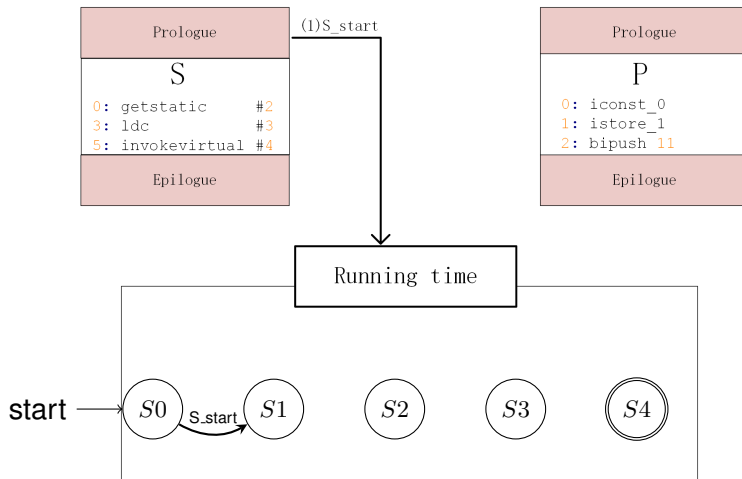
# Trigger



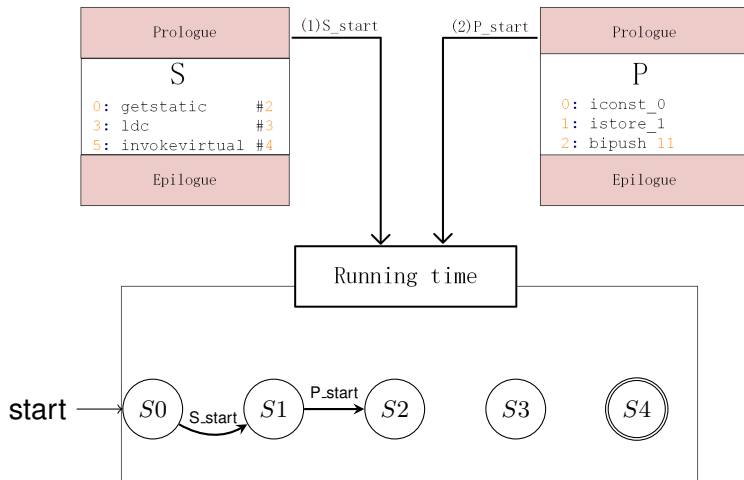
# Trigger



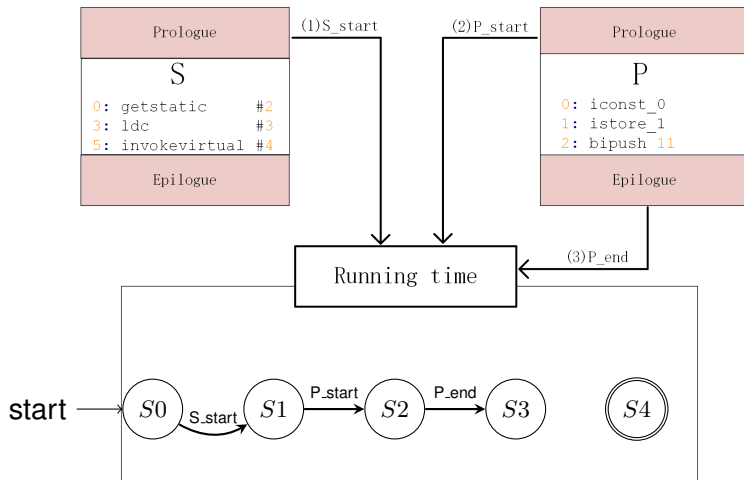
# Trigger



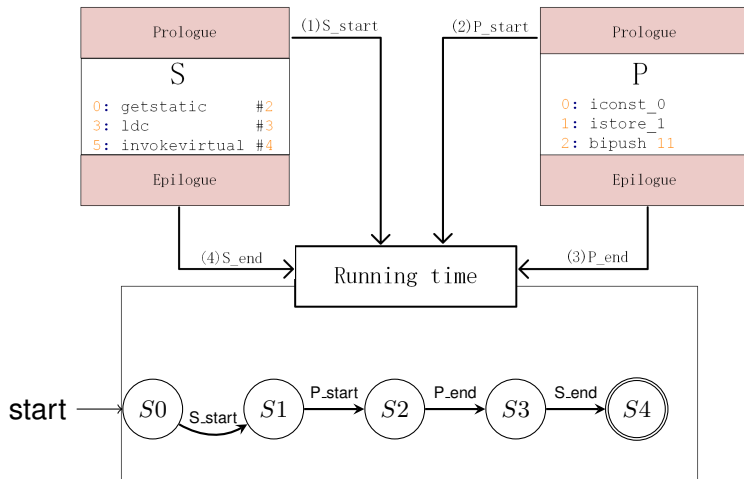
# Trigger



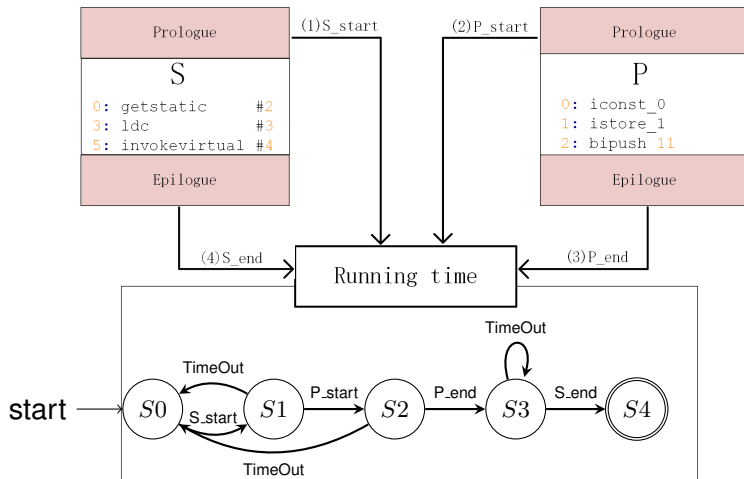
# Trigger



# Trigger



# Trigger



# Evaluation

Table: Systems under testing.

System	# CloudRaid code changes	Workload
Hadoop2/Yarn	48	wordcount + kill
HDFS	18	putfile + reboot
HBase	25	write + node crash
Cassandra	17	write



# Evaluation

Table: Message orders pruned by each analysis.

System	#Total	HB	% of Pruned		All
			Order	ID	
Hadoop2/Yarn	4489	1.0%	11.1%	81.5%	93.6%
HDFS	81	2.5%	45.7%	51.9%	85.2%
HBase	324	2.5%	57.7%	34.3%	94.4%
Cassandra	16	0.0%	75.0%	0.0%	75%

# Evaluation

**Table:** Analysis and testing times of CloudRaid.

System	Profiling(s)	Analysis(s)	Trigger(s)
Hadoop2/Yarn	648.0	131.3	6990.2
HDFS	646.0	60.0	828.3
HBase	1309.0	63.3	1368.0
Cassandra	263.1	112.3	60.3

# Evaluation

**Table:** 28 bugs detected and 8 of them are new bugs. All bugs are confirmed by the original developers, and 3 of them are already fixed.


Bug ID	type	status	Patched?	Symptom
YARN-6948	Order	Fixed	yes	Attempt fail
YARN-6969	Order	Unresolved	no	Wrong state
YARN-7176	Atomicity	Unresolved	yes	Cluster down
YARN-7563	Order	Unresolved	yes	Resource leak
YARN-7663	Order	Fixed	yes	Job fail
YARN-7726	Order	Unresolved	yes	Wrong state
YARN-7786	Order	Fixed	yes	Null Pointer
HBase-19004	Order	Unresolved	no	Data loss

# Why CloudRaid works

# Why CloudRaid works



# Why CloudRaid works



Distributed systems always  
log important messages

# Why CloudRaid works

Stop  
Task




Distributed systems always  
log important messages

# Why CloudRaid works




Distributed systems always  
log important messages







# YARN-7786




# YARN-7786




# YARN-7786




# YARN-7786




# YARN-7786



# YARN-7786




# YARN-7786






# YARN-7786



# YARN-7786



# YARN-7176



# YARN-7176


(1) startContainer




# YARN-7176




# YARN-7176



# YARN-7176




# YARN-7176







# YARN-7176




# YARN-7176





# YARN-7176





# YARN-7176



# YARN-7176







## Related work

- Concurrency bug detection
  - Model checking<sup>4</sup>: State-space explosion
  - DCatch<sup>5</sup>: Data Race, memory only;
- Log analysis<sup>6</sup>
  - Diagnose or monitor;
- Study on distributed bugs<sup>7</sup>

---

<sup>4</sup>leesatapornwongsa2014samc.

<sup>5</sup>liu2017dcatch.

<sup>6</sup>xu2009detecting.

<sup>7</sup>gunawi3000bugs; leesatapornwongsa2016taxdc.



# Conclusions

- CloudRaid:a simple yet effective tool to detect distributed concurrency bugs!

# Conclusions

- CloudRaid: a simple yet effective tool to detect distributed concurrency bugs!
- Testing 40 versions of 4 different systems in 35 hours!

# Conclusions

- CloudRaid: a simple yet effective tool to detect distributed concurrency bugs!
- Testing 40 versions of 4 different systems in 35 hours!
- Found 28 bugs, including 8 new bugs!

Thank You

Q & A