





计算技术研究所

优秀博士生论坛

 演讲人：陆杰

 指导老师：李炼（研究员，百人计划）



大纲



个人介绍



小组介绍




研究



问答



大纲



个人介绍



小组介绍



研究



问答



个人介绍

- 2014级直博生，计算所体系结构国家重点实验室
- 研究领域：程序分析，分布式系统，日志分析

第一作者

SOSP 19, FSE 18

SANER 19

2 CCF-A, 1 CCF-B

• 所长特别奖（夏培肃奖）

• 三好学生标兵 2019

• 国家奖学金 2019



大纲



个人介绍



小组介绍



研究



问答



小组介绍

程序分析小组

李炼

新南威尔士
大学

甲骨文

中科院百人计划





小组研究方向

通过**程序分析**技术来帮助提高
软件系统的**可靠性和安全性**
(**缺陷和漏洞检测**)



小组研究方向

静态&动态分析检测系统：Wukong



小组研究方向

静态&动态分析检测系统：Wukong

深度缺陷和漏洞



小组研究方向

静态&动态分析检测系统：Wukong

深度缺陷和漏洞

C/C++, Java, Go
Android, 分布式



小组研究方向

静态&动态分析检测系统：Wukong

深度缺陷和漏洞

C/C++, Java, Go
Android, 分布式

SOSP, FSE,
ASE, CGO等

Best Paper



大纲



个人介绍



小组介绍



研究





问答



计算技术研究所

分布式系统崩溃恢复 缺陷检测

 演讲人：陆杰


 指导老师：李炼（研究员，百人计划）




计算技术研究所

分布式系统 崩溃 恢复

缺陷 检测

 演讲人：陆杰


 指导老师：李炼（研究员，百人计划）




计算技术研究所

分布式系统 崩溃 恢复

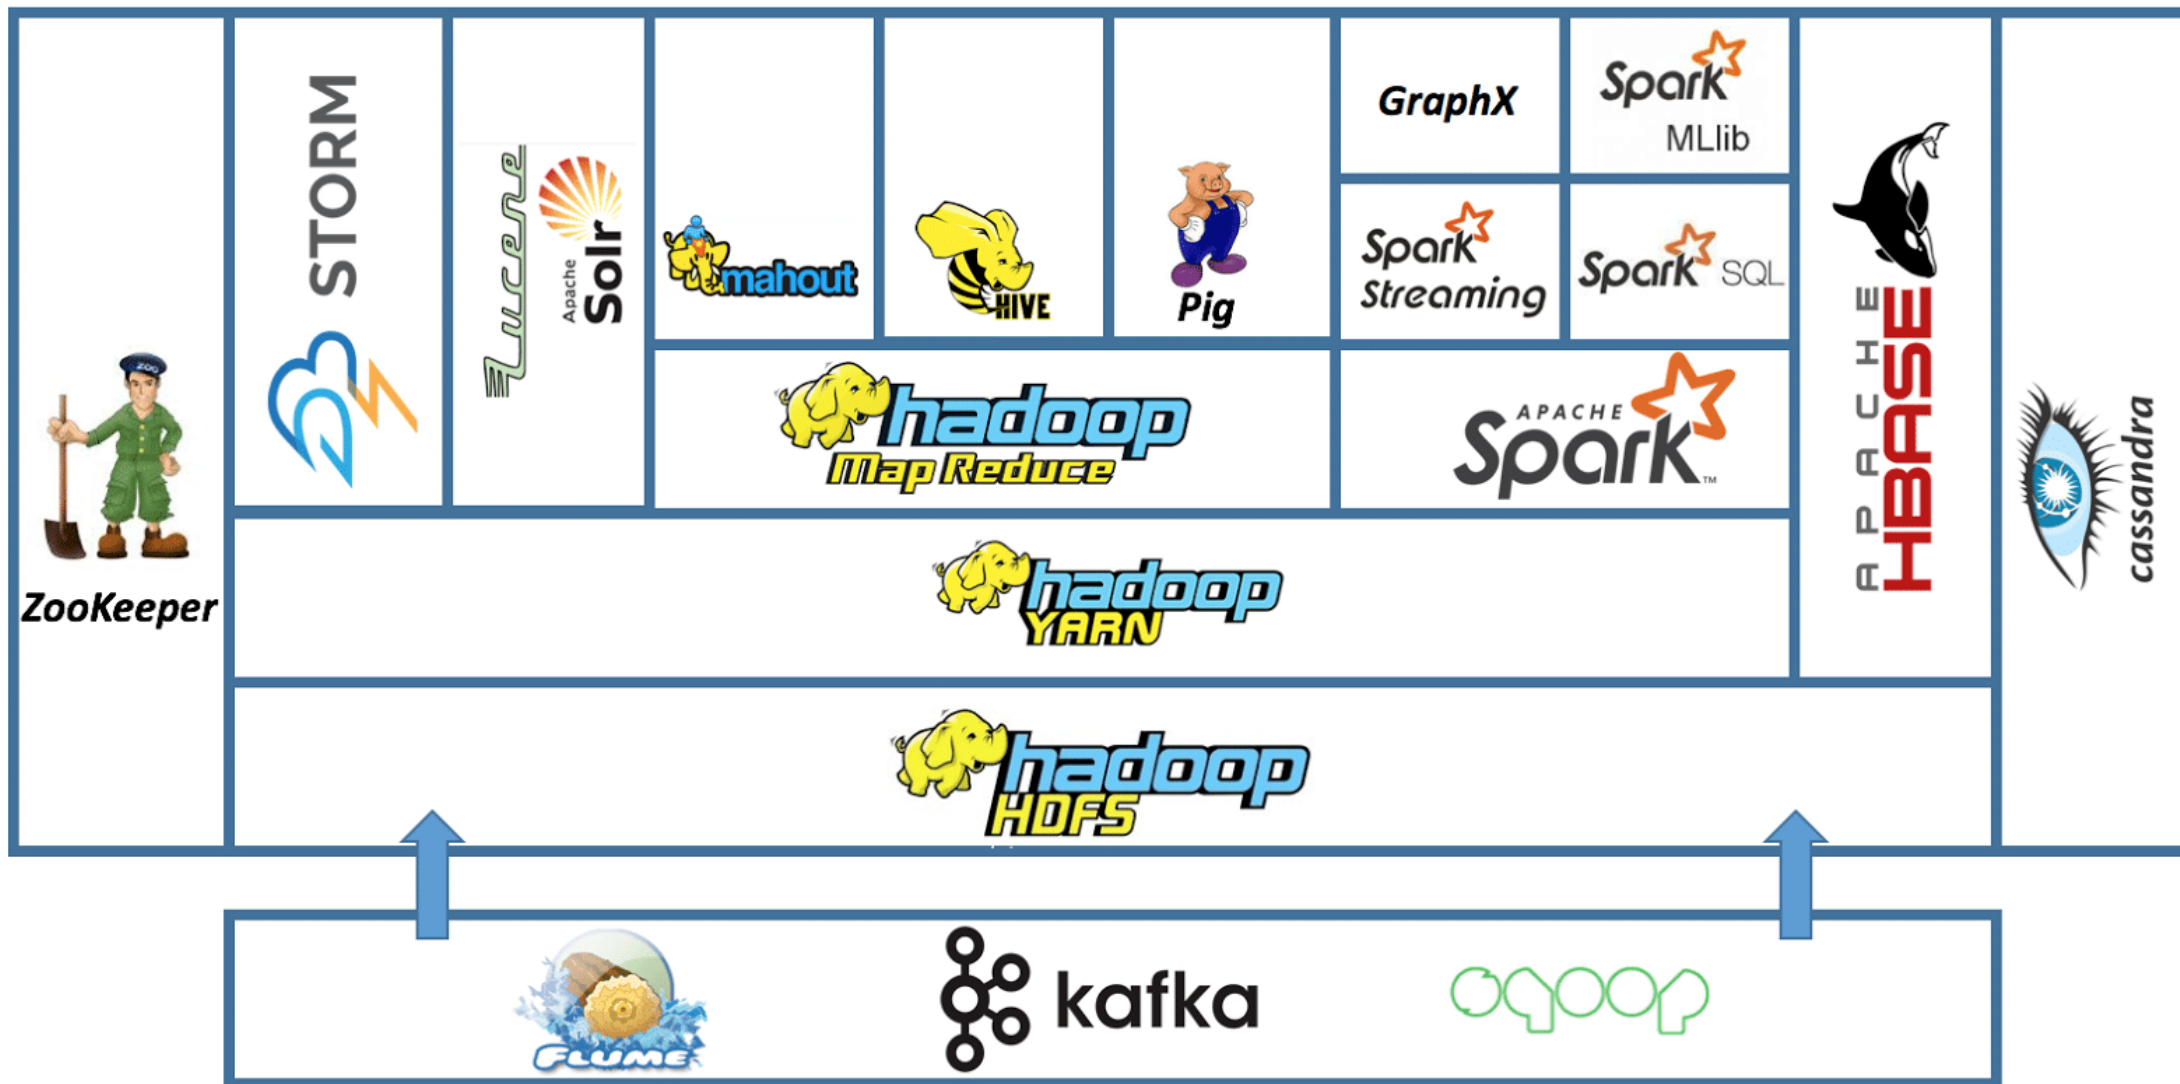
缺陷 检测

 演讲人：陆杰

 指导老师：李炼（研究员，百人计划）



分布式系统





分布式系统

Industry Cloud Solutions

Financial Services		Healthcare & Life Sciences		Manufacturing		Government		Energy		Retail	
Tech Vendors 	Enterprises 	Tech Vendors 	Enterprises 	Tech Vendors 	Enterprises 	Tech Vendors 	Agencies 	Tech Vendors 	Enterprises 	Tech Vendors 	Enterprises

Industry Clouds

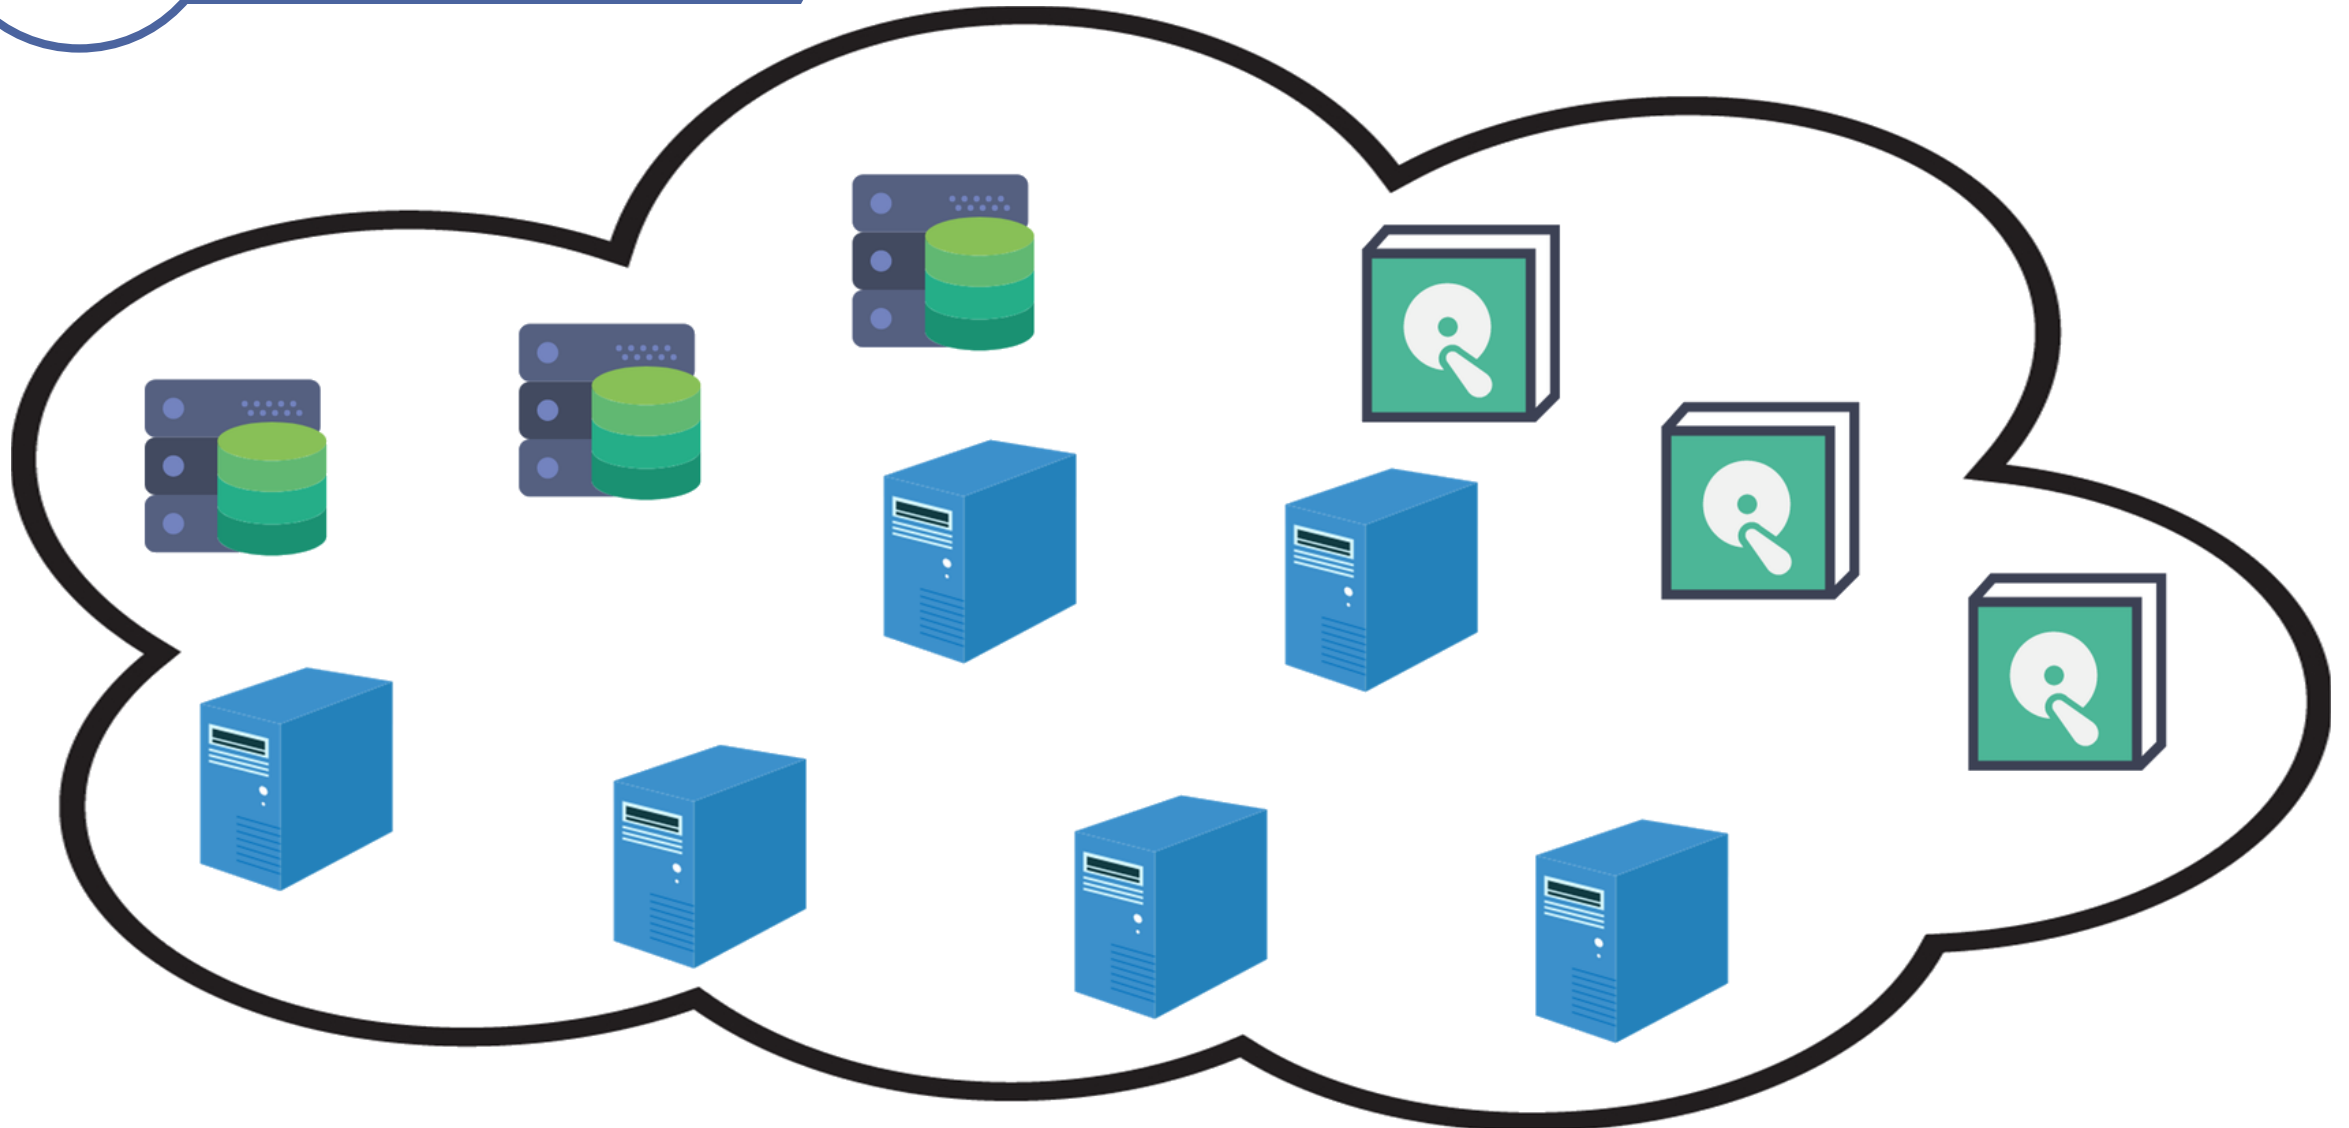
Information-Driven	Technology-Driven	Operations-Driven

Supporting Technology and Services

IaaS	PaaS	Professional Services



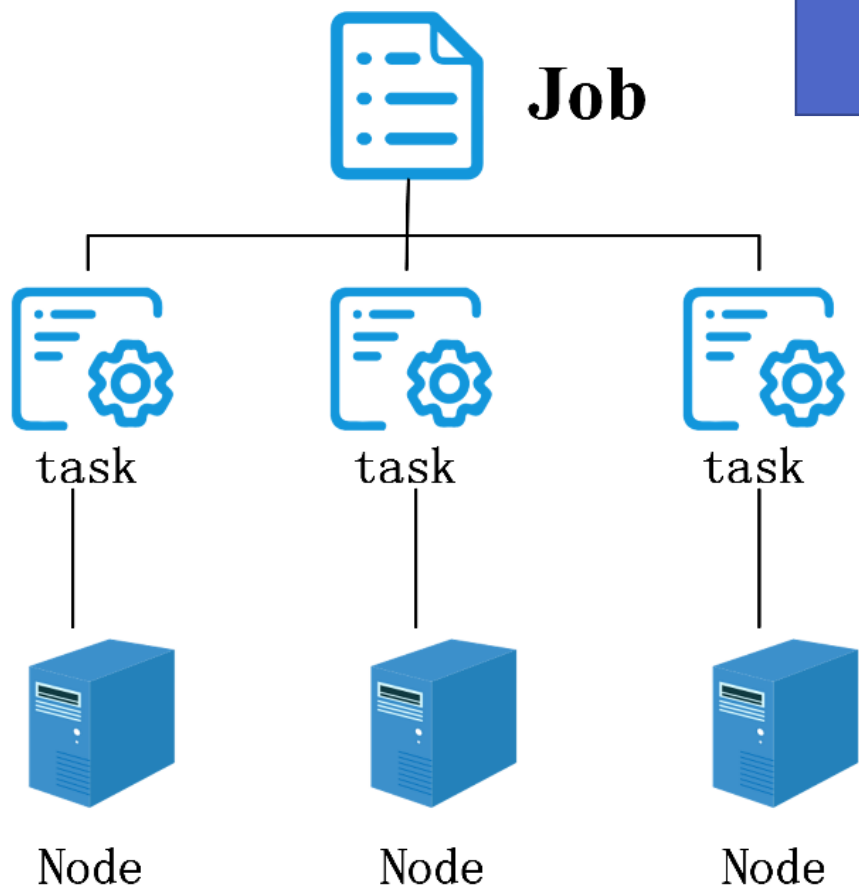
分布式系统





分布式系统

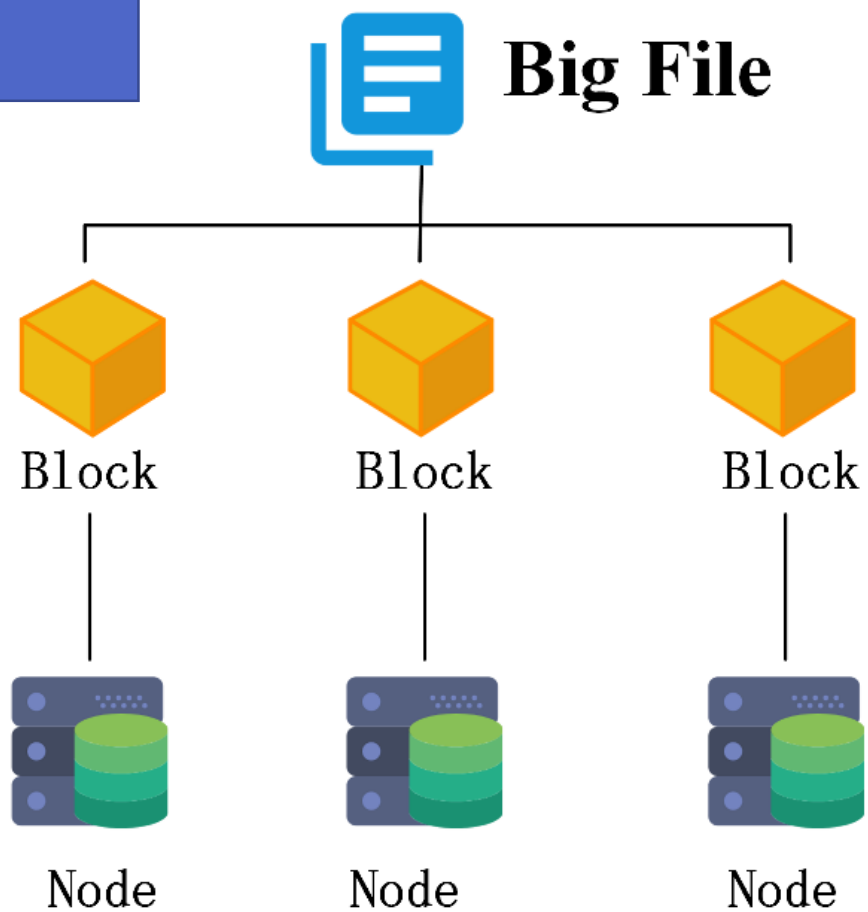

Scale





分布式系统

Scale







计算技术研究所

分布式系统 崩溃 恢复


缺陷 检测

 演讲人：陆杰

 指导老师：李炼（研究员，百人计划）




崩溃(Crash)





崩溃(crash)

每天至少180节点崩溃 (谷歌)







计算技术研究所

分布式系统 崩溃 恢复

缺陷 检测

 演讲人：陆杰

 指导老师：李炼（研究员，百人计划）



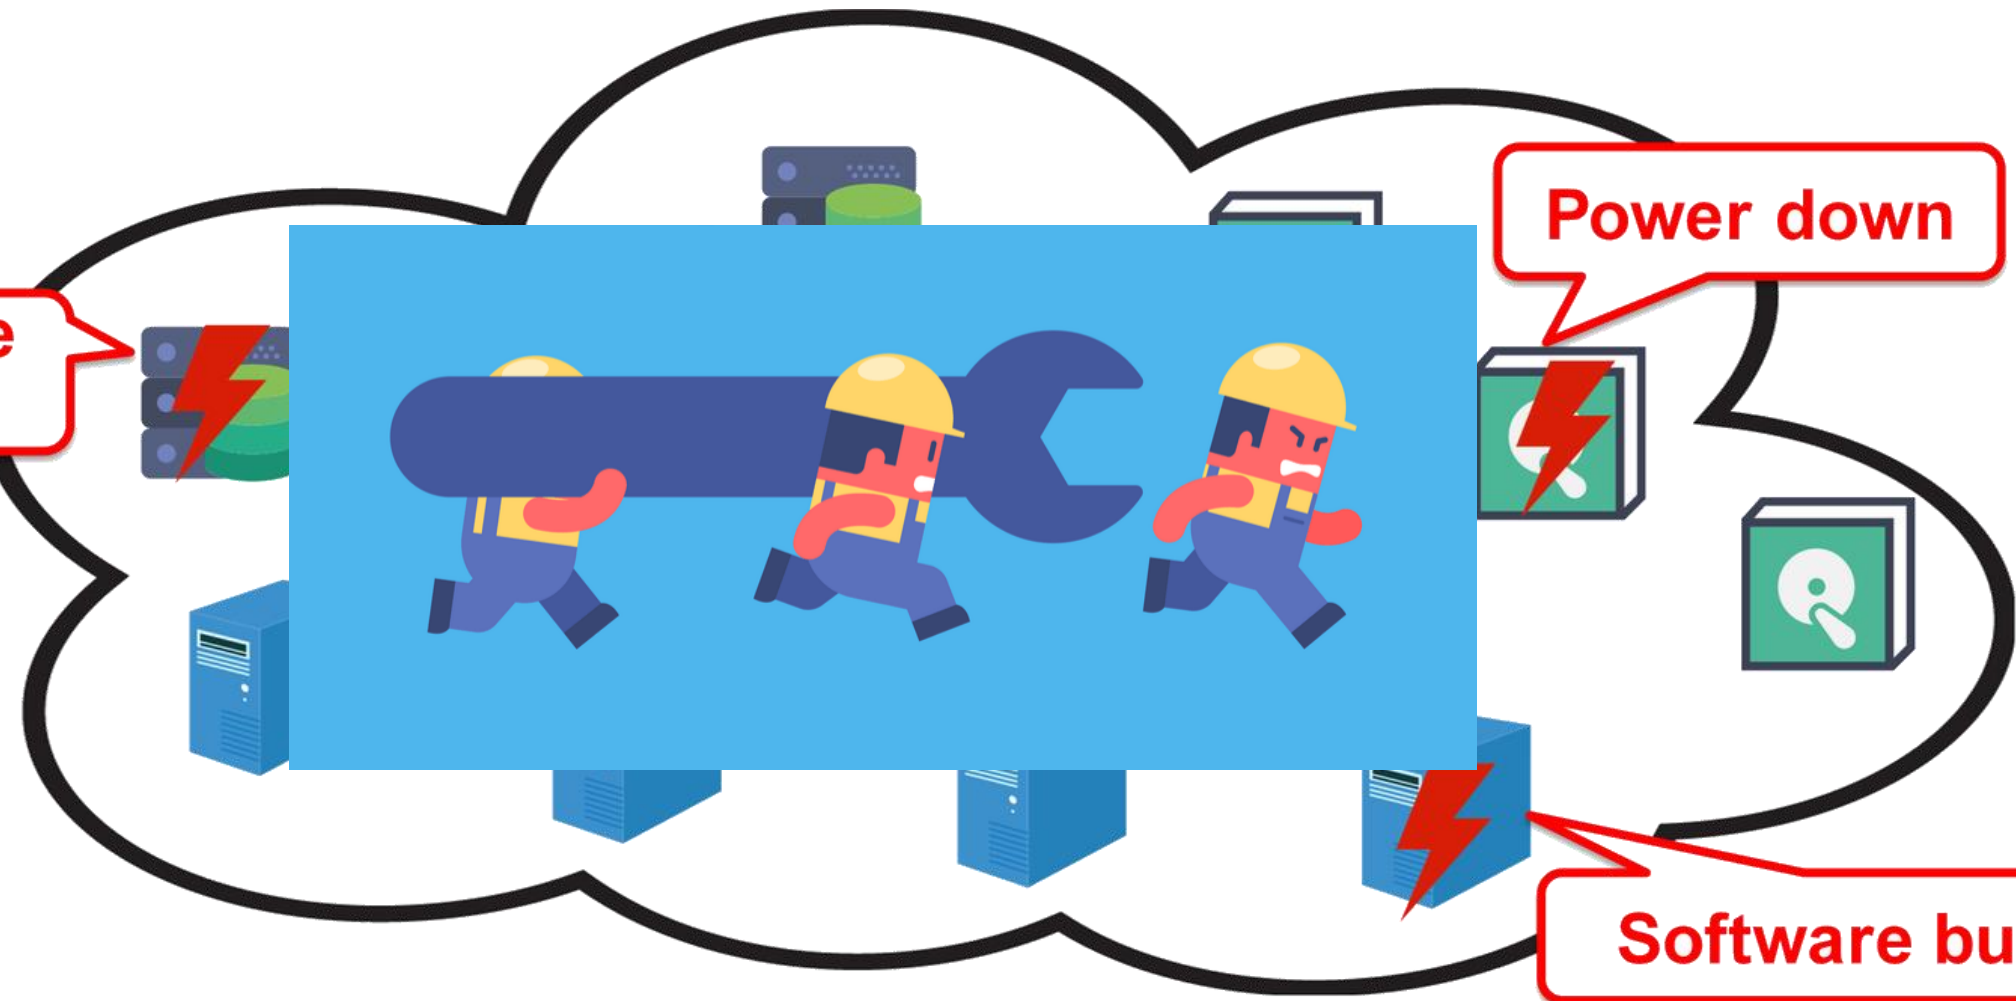
恢复(recovery)

恢复是分布式系统的
首要操作

Hardware
failure

Power down

Software bugs





恢复(recovery)

Job



task1.0



节点1

容错



恢复(recovery)

Job



task1.0



节点1

容错



恢复(recovery)

Job



task1.0



节点1

容错



恢复(recovery)



task1.1

Job



task1.0



节点2



节点1

容错



恢复(recovery)



task1.1

Job



task1.0



节点2




节点1

更新
元信息变量



元信息变量



Cluster






元信息变量


Cluster




node0



node1




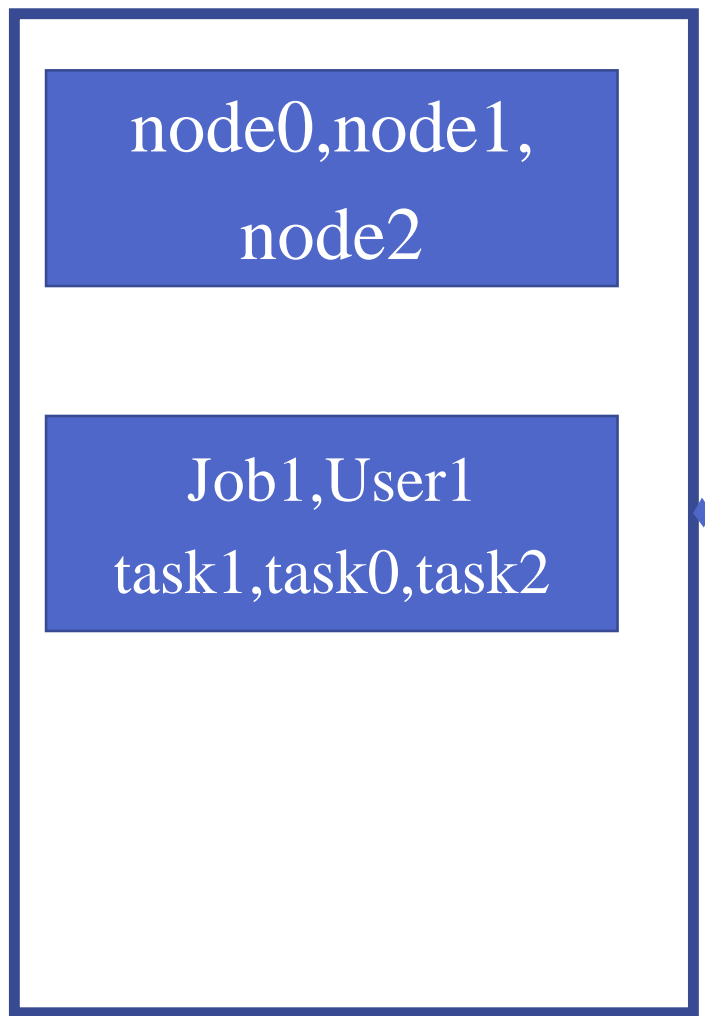
node2






元信息变量


Cluster



node0
task2



node1
task1


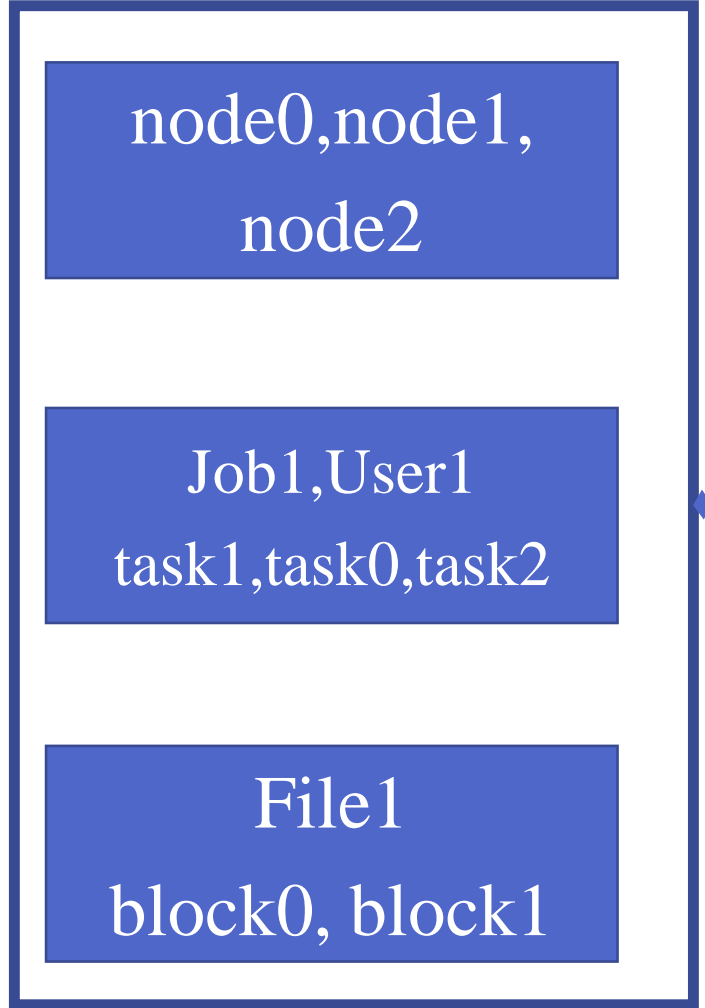


node2
task0




元信息变量

Cluster



node0
task2
block1



node1
task1
block0




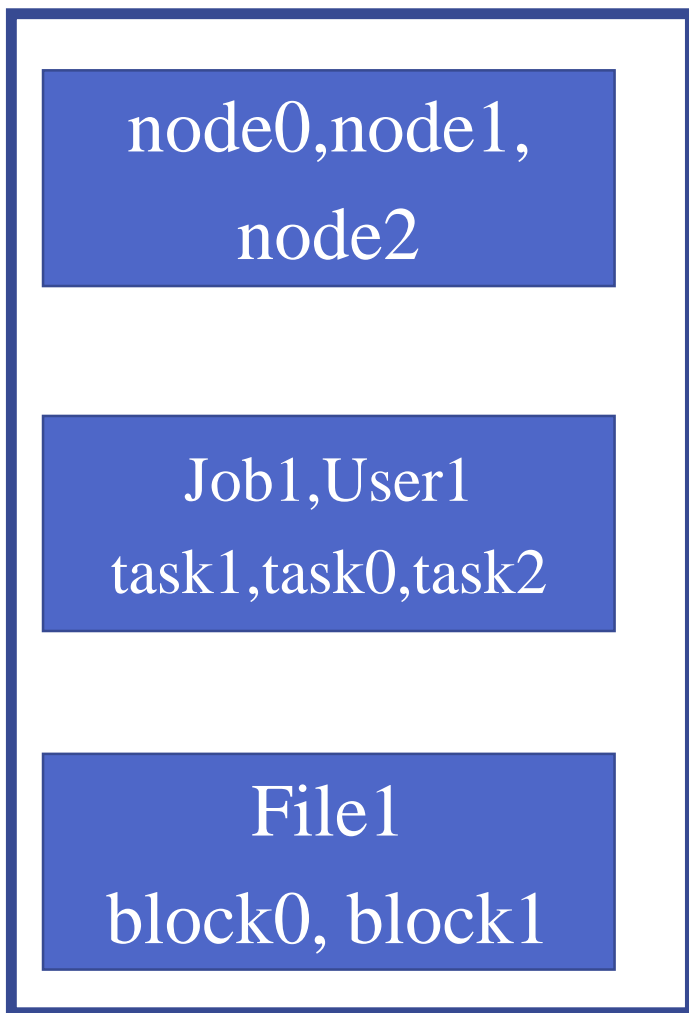
node2
task0




元信息变量

元信息:nodes, tasks,job,block等


Cluster



node0
task2
block1



node1
task1
block0



node2
task0




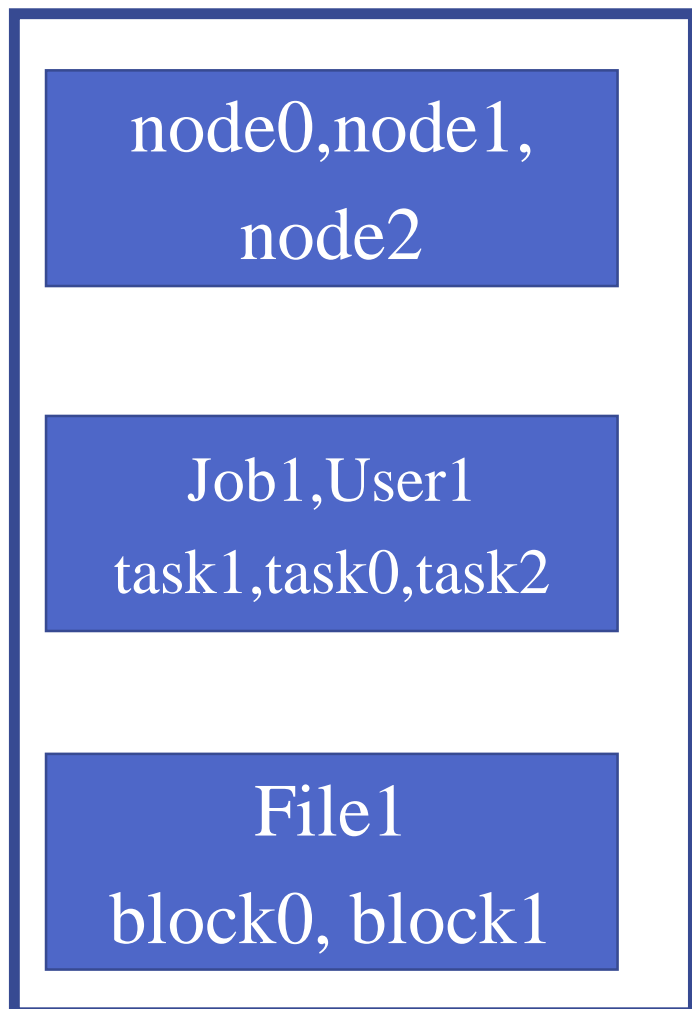
元信息变量

元信息:nodes, tasks,job,block等

元信息变量



Cluster





元信息变量

元信息:nodes, tasks,job,block等

Cluster

- livingNodes
- livingTasks
- blocks

node0,node1,
node2

Job1,User1
task1,task0,task2

File1
block0,block1



node0
task2
block1



node1
task1
block0



node2
task0

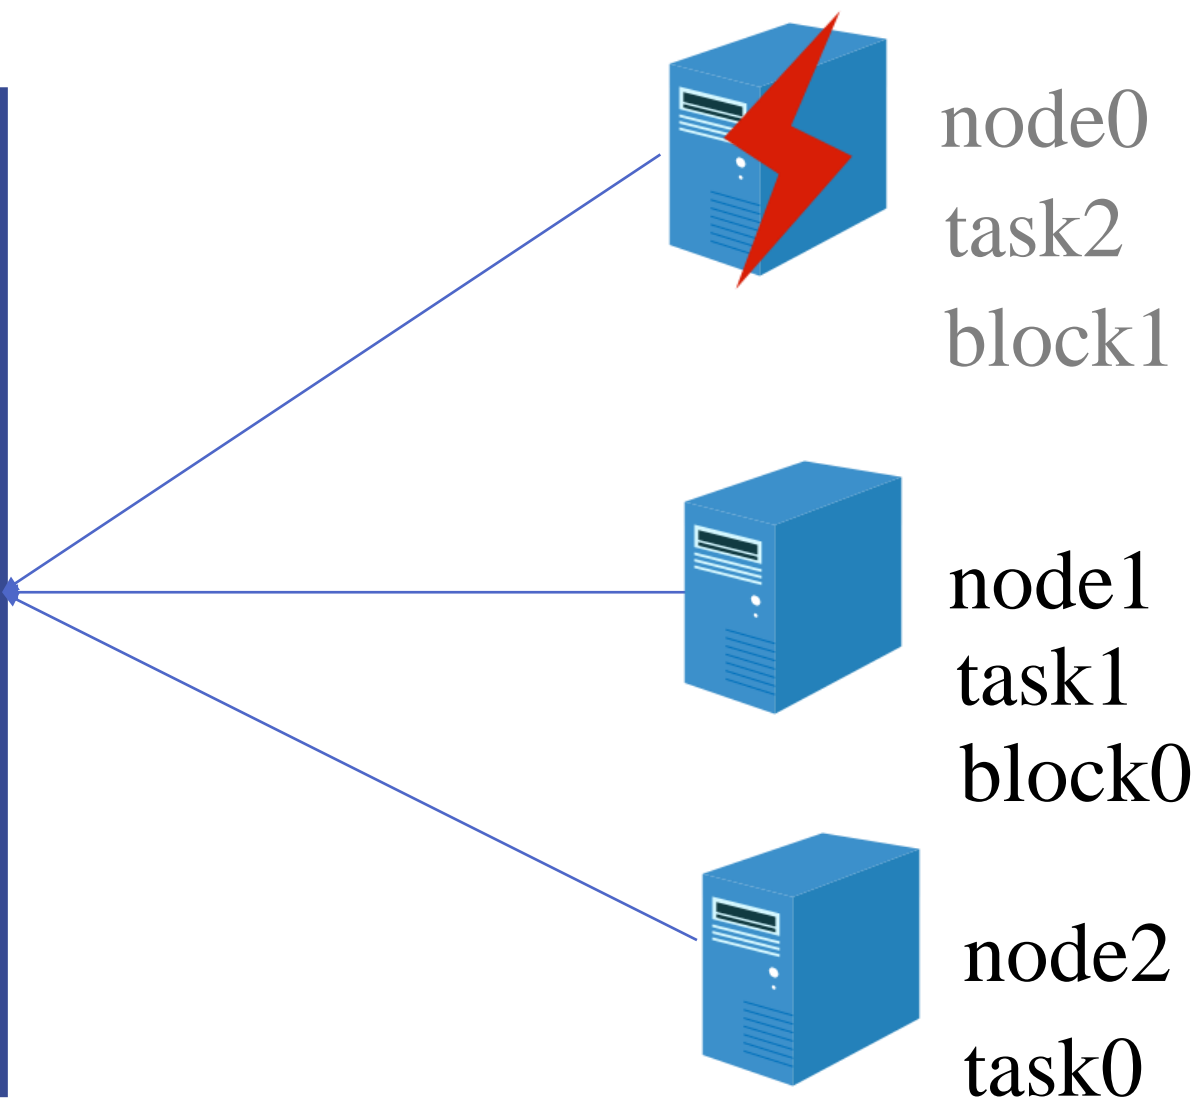


元信息变量

元信息:nodes, tasks,job,block等


元信息变量

Cluster



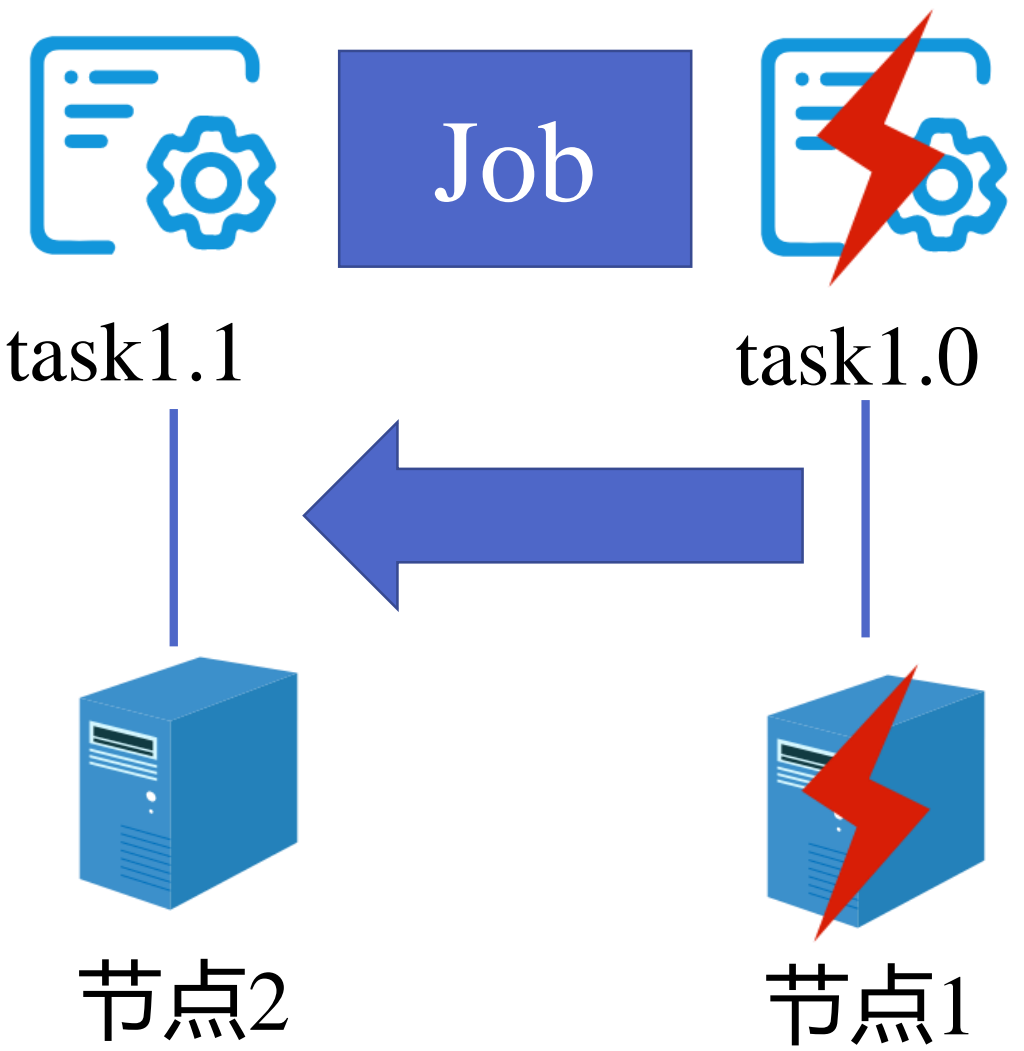


小结






小结






小结



恢复操作



小结



恢复操作




更新元信息变量




计算技术研究所

分布式系统 崩溃 恢复

缺陷 检测

 演讲人：陆杰

 指导老师：李炼（研究员，百人计划）

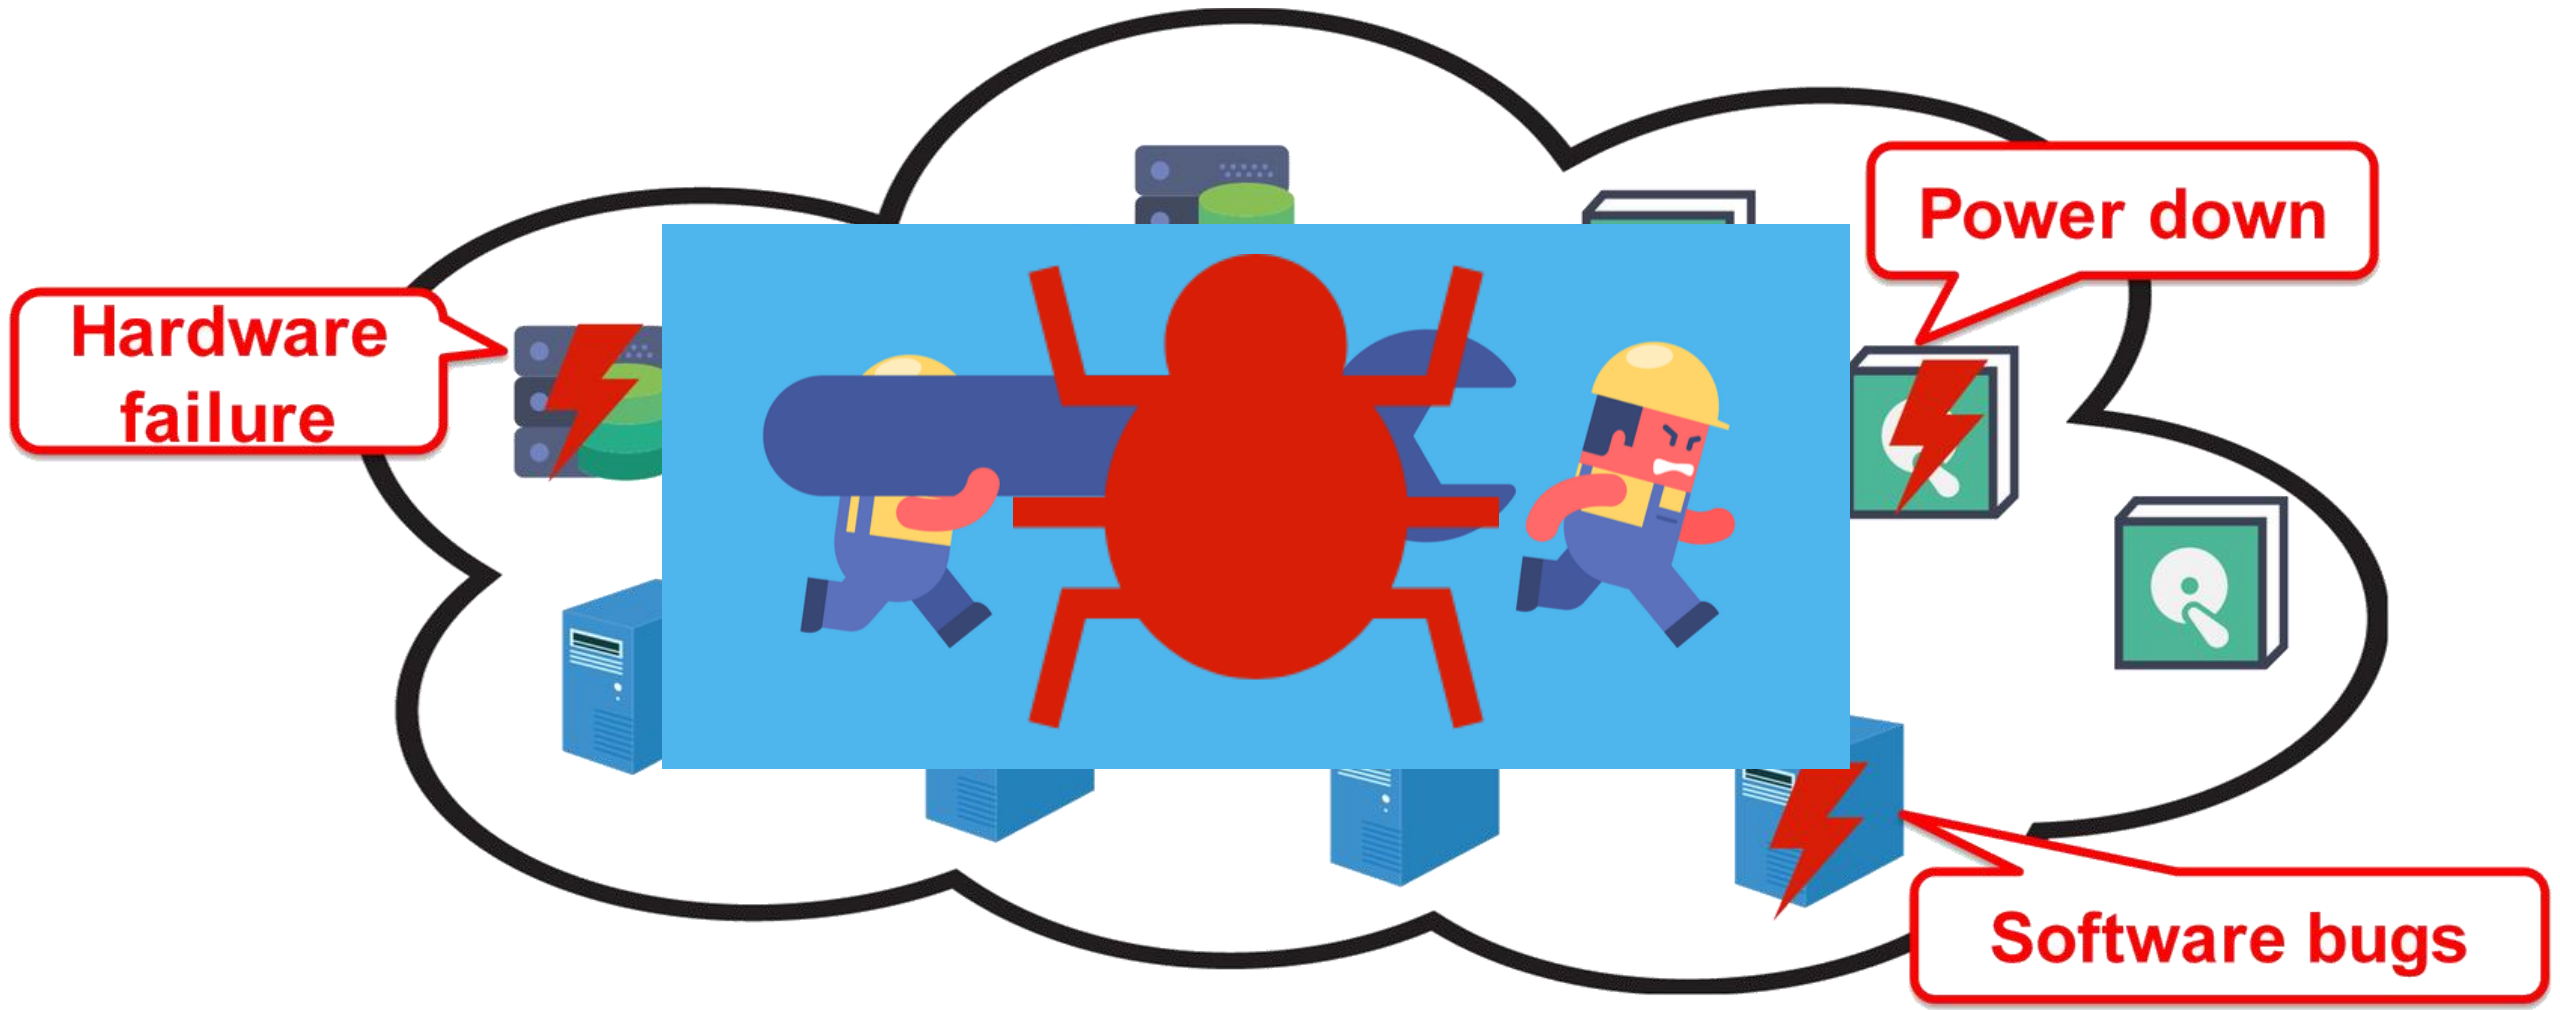


崩溃恢复缺陷





崩溃恢复缺陷





崩溃恢复缺陷

Cluster



node2





崩溃恢复缺陷

1. NodeId nodeId = “node2”
2. livingNodes.add(nodeId);

Cluster



node2





崩溃恢复缺陷

Cluster

node2

1. NodeId nodeId = "node2"
2. livingNodes.add(nodeId);





崩溃恢复缺陷

Cluster

node2

1. NodeId nodeId = "node2"
2. livingNodes.add(nodeId);



livingNodes.remove(nodeId);

恢复



崩溃恢复缺陷

Cluster


node2

1. NodeId nodeId = "node2"
2. livingNodes.add(nodeId);



livingNodes.remove(nodeId);

恢复






崩溃恢复缺陷

Cluster

node2

```
1. NodeId nodeId = "node2"  
2. livingNodes.add(nodeId);  
←  
while(livingNodes.has("node2")) {  
    retryConnet("node2");  
}  
livingNodes.remove(nodeId);
```

恢复






崩溃恢复缺陷

Cluster


node2

1. NodeId nodeId = "node2"
2. livingNodes.add(nodeId);

```
while(livingNodes.has("node2")) {  
    retryConnet("node2");  
}  
livingNodes.remove(nodeId);
```



恢复






崩溃恢复缺陷

写之后(post-write)


```
1. NodeId nodeId = "node2"  
2. livingNodes.add(nodeId);  
←  
while(livingNodes.has("node2")) {  
    retryConnet("node2");  
}  
livingNodes.remove(nodeId);
```

Cluster

node2



恢复





崩溃恢复缺陷

```
int length = livingNodes.size();  
(资源调度)
```

```
livingNodes.get(length-1);
```

Cluster



node2





崩溃恢复缺陷

Cluster

node2

```
int length = livingNodes.size();  
(资源调度)
```

```
livingNodes.get(length-1);
```





崩溃恢复缺陷

Cluster

node2

```
int length = livingNodes.size();  
(资源调度)
```

```
livingNodes.get(length-1);
```

```
livingNodes.remove(nodeId)
```

恢复





崩溃恢复缺陷

Cluster

node2

```
int length = livingNodes.size();  
(资源调度)
```

```
livingNodes.get(length-1);
```

```
livingNodes.remove(nodeId)
```



恢复





崩溃恢复缺陷

读之前(pre-read)


```
int length = livingNodes.size();
```

(资源调度)

```
livingNodes.get(length-1);
```

```
livingNodes.remove(nodeId)
```

Cluster



恢复

node2





崩溃恢复缺陷



崩溃恢复缺陷

崩溃恢复缺陷两种模式：

元信息变量的


读之前，写之后




计算技术研究所

分布式系统 崩溃 恢复

缺陷 检测

 演讲人：陆杰

 指导老师：李炼（研究员，百人计划）



检测设计

1. 识别所有的元信息变量



检测设计

1. 识别所有的元信息变量

元信息变量

1. NodeId nodeId = "node2"
2. **livingList**.add(nodeId);

写或读

node1

2. 崩溃点



检测设计

1. 识别所有的元信息变量

元信息变量

1. NodeId nodeId = "node2"
2. **livingList**.add(nodeId);

写或读

node1



2. 崩溃点



node2



3





检测设计

1. 识别所有的元信息变量

元信息变量

1. NodeId nodeId = "node2"
2. **livingList**.add(nodeId);

写或读

.....



node1

2. 崩溃点

恢复

node2

3





识别元信息变量

节点变量是最基础的元信息变量

```
1 public void registerNode(NodeId nodeId) {  
2     LOG.info("registering node " + nodeId);  
3 }
```

10.5.1.11
节点的IP



识别元信息变量

和元信息变量**相关**的变量也是
元信息变量

```
1 public void assignTask(NodeId nodeId,  
2   ContainerId containerId) {  
3   LOG.info("Using container " + containerId +  
4           "on node" + nodeId);  
5 }
```

元信息变量

元信息变量



识别元信息变量

元信息

```
1 public void registerNode(NodeId nodeId) {  
2     Log.info("registering node" + nodeId);  
3     RPCfun(nodeId);  
4 }  
5 public void RPCfun(NodeId nodeId) {  
6     livingNodes.add(nodeId);  
7 }
```



识别元信息变量

元信息

```
1 public void registerNode(NodeId nodeId) {  
2     Log.info("registering node" + nodeId);  
3     RPCfun(nodeId);  
4 }  
5 public void RPCfun(NodeId nodeId) {  
6     livingNodes.add(nodeId);  
7 }
```



识别元信息变量

元信息

```
1 public void registerNode(NodeId nodeId) {  
2     Log.info("registering node" + nodeId);  
3     RPCfun(nodeId);  
4 }  
5 public void RPCfun(NodeId nodeId) {  
6     livingNodes.add(nodeId);  
7 }
```

指针分析



识别元信息变量

元信息

```
1 public void registerNode(NodeId nodeId) {  
2     Log.info("registering node" + nodeId);  
3     RPCfun(nodeId);  
4 }  
5 public void RPCfun(NodeId nodeId) {  
6     livingNodes.add(nodeId);  
7 }
```

指针分析

元信息变量



识别元信息变量

元信息

```
1 public void registerNode(NodeId nodeId) {  
2     Log.info("registering node" + nodeId);  
3     RPCfun(nodeId);  
4 }  
5 public void RPCfun(NodeId nodeId) {  
6     livingNodes.add(nodeId);  
7 }
```

元信息变量



识别元信息变量

元信息

```
1 public void registerNode(NodeId nodeId) {  
2     Log.info("registering node" + nodeId);  
3     RPCfun(nodeId);  
4 }  
5 public void RPCfun(NodeId nodeId) {  
6     livingNodes.add(nodeId);  
7 }
```

类型分析

元信息变量

```
List<NodeId> livingNodes = null;
```



细节

SOSP 2019

CrashTuner: Detecting Crash-Recovery Bugs in Cloud Systems via Meta-Info Analysis

Jie Lu

lujie@ict.ac.cn

SKL Computer Architecture, ICT, CAS
University of Chinese Academy of
Sciences, China

Chen Liu

liuchen17z@ict.ac.cn

SKL Computer Architecture, ICT, CAS
University of Chinese Academy of
Sciences, China

Lian Li[†]

lianli@ict.ac.cn

SKL Computer Architecture, ICT, CAS
University of Chinese Academy of
Sciences, China

Xiaobing Feng

xfb@ict.ac.cn

SKL Computer Architecture, ICT, CAS
University of Chinese Academy of
Sciences, China

Feng Tan

Jun Yang

Liang You

{tanfeng.tf,muzhuo.yj,youliang.yl}@alibaba.com
Alibaba Group

Abstract

Crash-recovery bugs (bugs in crash-recovery-related mechanisms) are among the most severe bugs in cloud systems and can easily cause system failures. It is notoriously difficult to detect crash-recovery bugs since these bugs can only be exposed when nodes crash under special timing conditions. This paper presents CrashTuner, a novel fault-injection testing approach to combat crash-recovery bugs. The novelty of CrashTuner lies in how we identify fault-injection points (crash points) that are likely to expose errors. We observe that if a node crashes while accessing *meta-info* variables, i.e., variables referencing high-level system state information (e.g., an instance of node or task), it often triggers crash-recovery bugs. Hence, we identify crash points by automatically inferring meta-info variables via a log-based static program analysis. Our approach is automatic and no manual specification is required.

We have applied CrashTuner to five representative distributed systems: Hadoop2/Yarn, HBase, HDFS, ZooKeeper, and Cassandra. CrashTuner can finish testing each system in 17.39 hours, and reports 21 new bugs that have never been found before. All new bugs are confirmed by the original developers and 16 of them have already been fixed (14 with

our patches). These new bugs can cause severe damages such as cluster down or start-up failures.

CCS Concepts • Software and its engineering → Software testing and debugging; Cloud computing.

Keywords Crash Recovery Bugs; Fault Tolerance; Distributed Systems; Bug Detection; Fault Injection; Cloud Computing

ACM Reference Format:

Jie Lu, Chen Liu, Lian Li, Xiaobing Feng, Feng Tan, Jun Yang, and Liang You. 2019. CrashTuner: Detecting Crash-Recovery Bugs in Cloud Systems via Meta-Info Analysis. In *ACM SIGOPS 27th Symposium on Operating Systems Principles (SOSP '19)*, October 27–30, 2019, Huntsville, ON, Canada. ACM, New York, NY, USA, 17 pages. <https://doi.org/10.1145/3341301.3359645>

1 Introduction

Distributed systems have become the backbone of computing in the cloud era. More and more applications are built on top of large-scale distributed systems (such as scalable computing frameworks [20, 57] and distributed storage systems [22, 36]), to provide online services to users. High availability of those systems is crucial: failures of the underlying distributed systems can lead to cloud outage, easily costing service providers millions of dollars [2, 9].

High availability of distributed systems largely hinges on how well these systems tolerate node crashes (failures). Large-scale distributed systems are often comprised of thousands of nodes (machines) [55], and it is common that a node may fail due to hardware or software faults [49]. Although various sophisticated crash-recovery mechanisms [4, 13, 16] have been adopted in distributed systems, it is still challenging to handle node crashes correctly. It is very difficult, if not impossible, for developers to anticipate all possible crash scenarios and correctly implement corresponding recovery mechanisms. In this paper, we refer to bugs in crash-recovery-related mechanisms as *crash-recovery bugs*.

^{*}corresponding author: lianli@ict.ac.cn

[†]Also with TianQi Soft Inc., China.

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than ACM must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

SOSP '19, October 27–30, 2019, Huntsville, ON, Canada

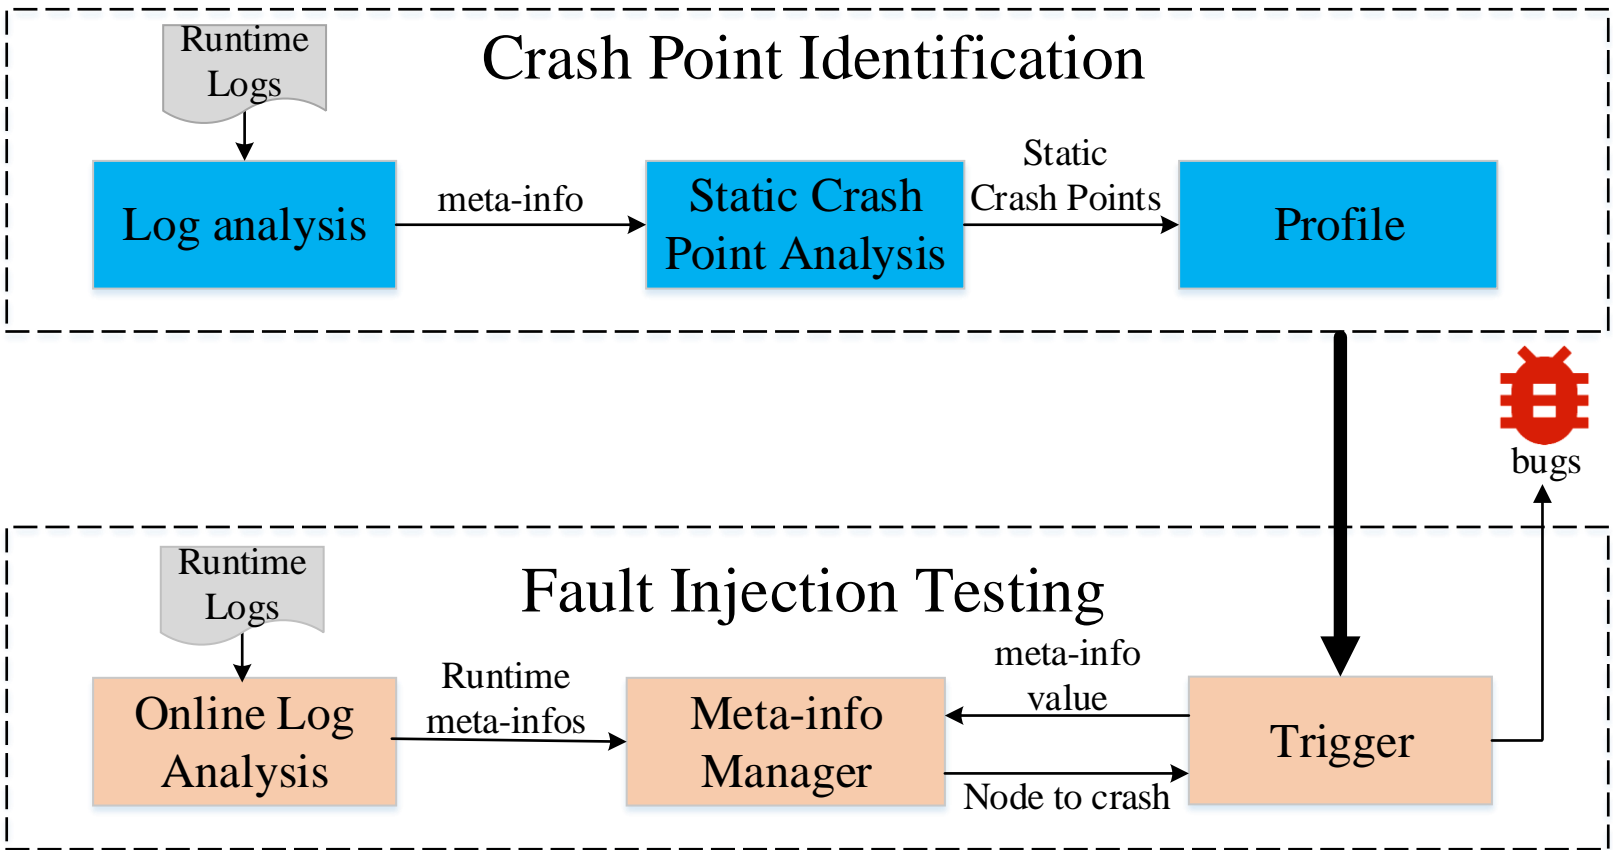
© 2019 Association for Computing Machinery.

ACM ISBN 978-1-4503-6873-5/19/10...\$15.00

<https://doi.org/10.1145/3341301.3359645>



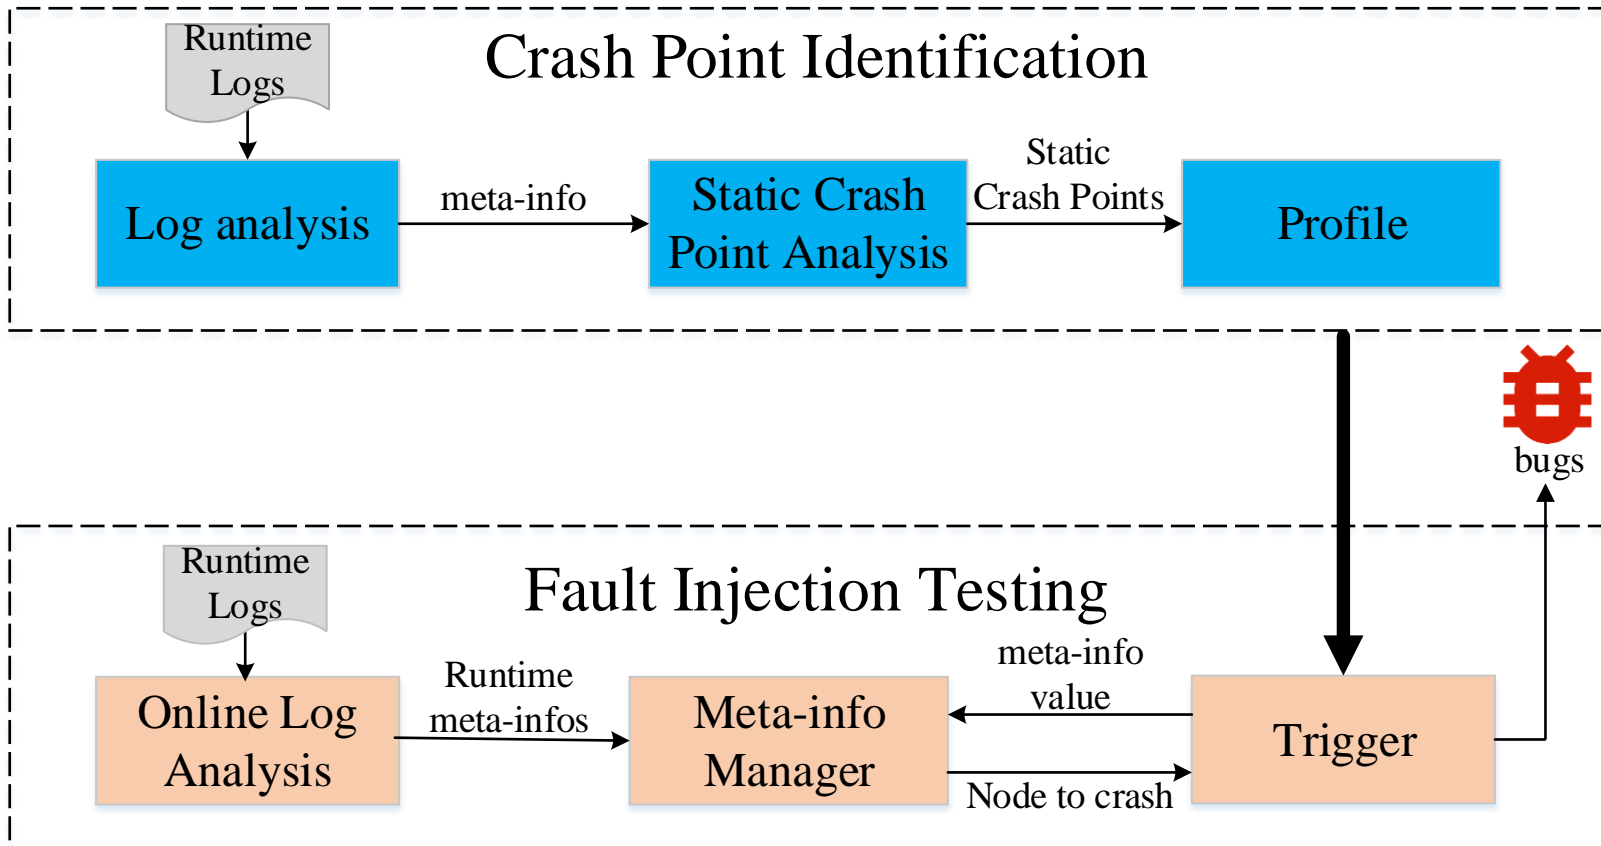
CrashTuner





CrashTuner

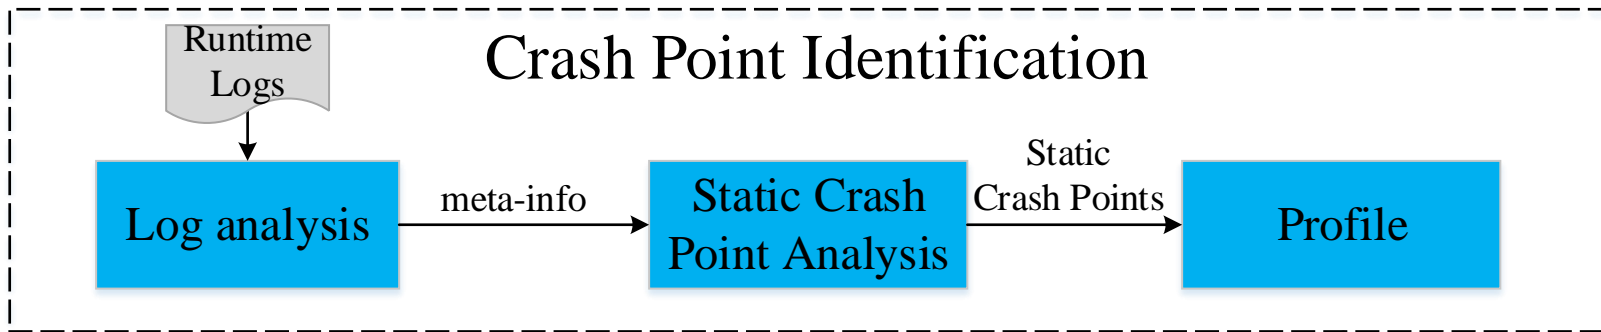
YARN, HDFS, Cassandra, Hbase, Zookeeper



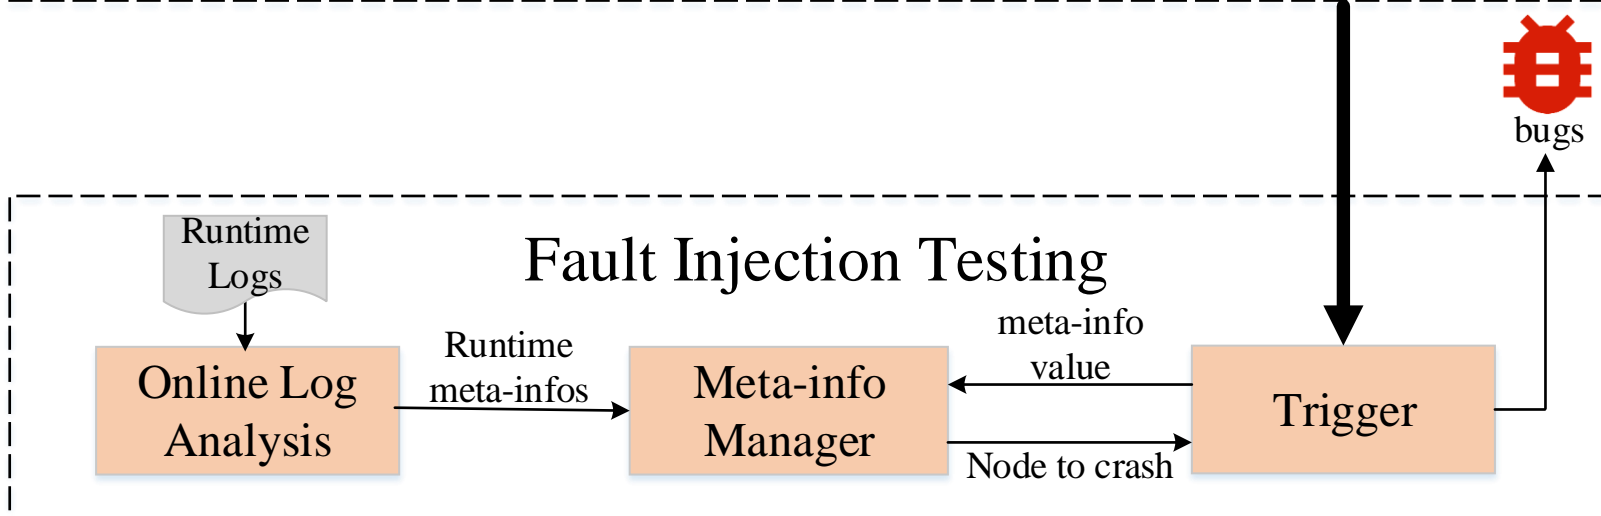


CrashTuner

YARN, HDFS, Cassandra,
Hbase, Zookeeper




17个小时内



21个缺陷
10个critical
16被修复

展示完毕 感谢各位聆听

 答辩人：陆杰

 指导教师：李炼

